

# 113: Architecture

Spring 2018

*Lecture:* Machine-Level Programming I

**Instructor:** Dr. Jana Giceva

# Textbooks

- ***“Computer systems: A programmer’s perspective”***  
Randal E. Bryant and David O’Hallaron, 2013 (3<sup>rd</sup> edition)

Available as e-book in the Imperial library

Website: <http://csapp.cs.cmu.edu>


Very important for the course:  
practice problems relevant to exam questions

- ***Intel 64 and IA-32 manuals*** for extra reference
- Not mandatory: any good C book

# Hardware/software interface

- What is software? What is hardware? Why do we need HW/SW interface?
- Why do we need to understand both sides of this interface?

```
3 public class DebuggingDemo {  
4  
5     int counter;  
6  
7     public void increment()  
8     {  
9         int temp = counter;  
10        temp ++;  
11        counter = temp;  
12    }  
13  
14    public static void main(String[] args) {  
15        DebuggingDemo d = new DebuggingDemo();  
16        d.setCounter(1);  
17        d.increment();  
18        System.out.println(d.counter);  
19    }  
20  
21    public void setCounter(int p_initialValue)  
22    {  
23        counter = p_initialValue;  
24    }  
25 }
```



# C/Java, assembly, and machine code

C / Java

```
if (x != 0) {  
    y = (y+z)/x;  
}
```

Assembly

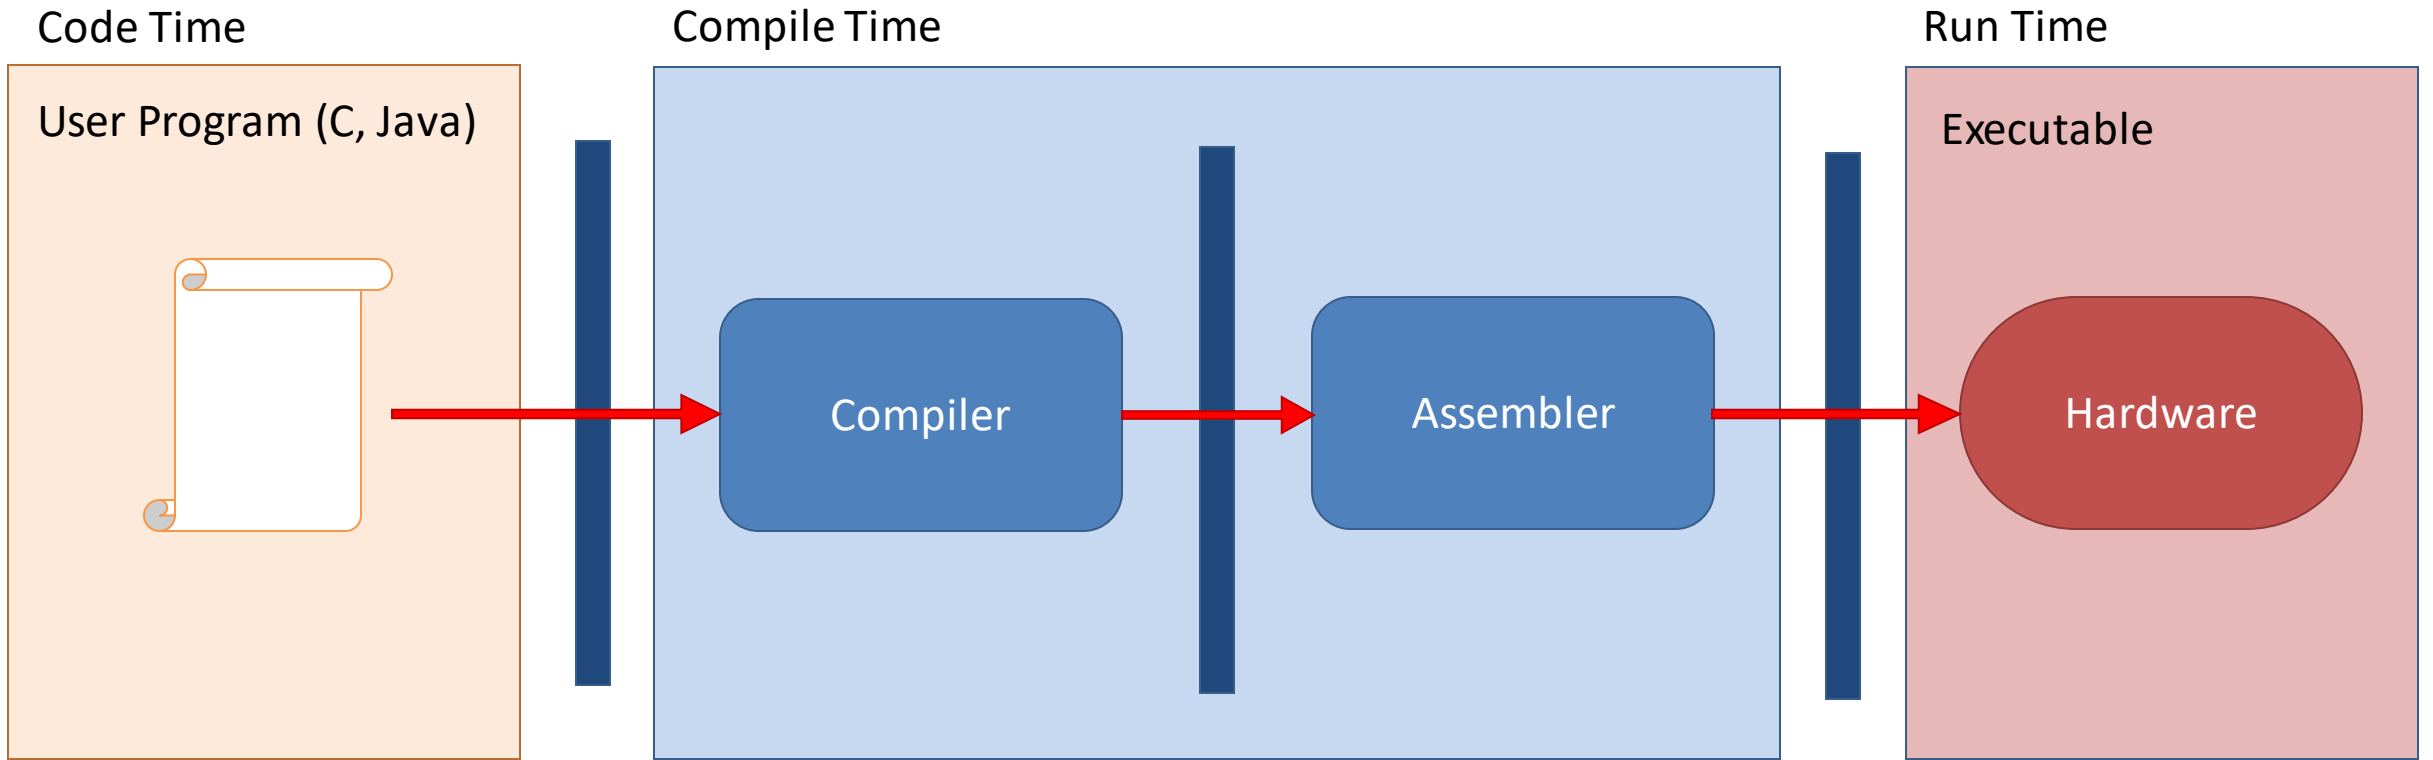
```
    cmpl    $0,-4(%rbp)  
    je      .L2  
    movl    -8(%rbp), %edx  
    movl    -12(%rbp), %eax  
    addl    %edx, %eax  
    cltd  
    idivl   -4(%rbp)  
    move    %eax, -8(%rbp)  
.L2:
```

Machine code

```
10000011011111000010010000011100000000  
0 0111010000011000  
10001011010001000010010000010100  
10001011010001100010010100010100  
1000110100000100000000010  
1000100111000010  
110000011111101000011111  
11110111011111000010010000011100  
10001001010001000010010000011000
```

- All program fragments are equivalent
- You'd rather write C / Java (more human friendly)
- Hardware executes strings of bytes

# Code / Compiler / Run time



- Note: The compiler and assembler are just programs (developed using the same principle).

# Roadmap – Overview of material

- Memory and data
- Machine code and x86 assembly language
- Procedures and stacks
- Arrays, structs and objects
- Floating point numbers and representation
- Memory architecture and caches
- Processes, interrupts and exceptions

# Writing assembly code? In 2018??

- You are probably never going to write a program in assembly
  - compilers are much better and more patient
  - unless you write a delicate, “special” code
- But, understanding assembly is **key** to the machine level execution model
  - behaviour of programs in the presence of **bugs**
    - High-level language model breaks down
  - tuning program **performance**
    - understand the optimisations done / not done by the compiler
    - understanding sources of program inefficiency
  - implementing **system software** (e.g., Operating Systems, Compilers)
  - creating / fighting **malware**

# Memory matters!

- RAM is an unrealistic abstraction:
  - Memory is not unbounded
  - Memory performance is not uniform (caches, virtual memory, etc.)
  - Memory referencing bugs are especially tricky (working with pointers in C).

```
void copy_ij(int src[2048][2048],
             int dst[2048][2048]) {
    int i, j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

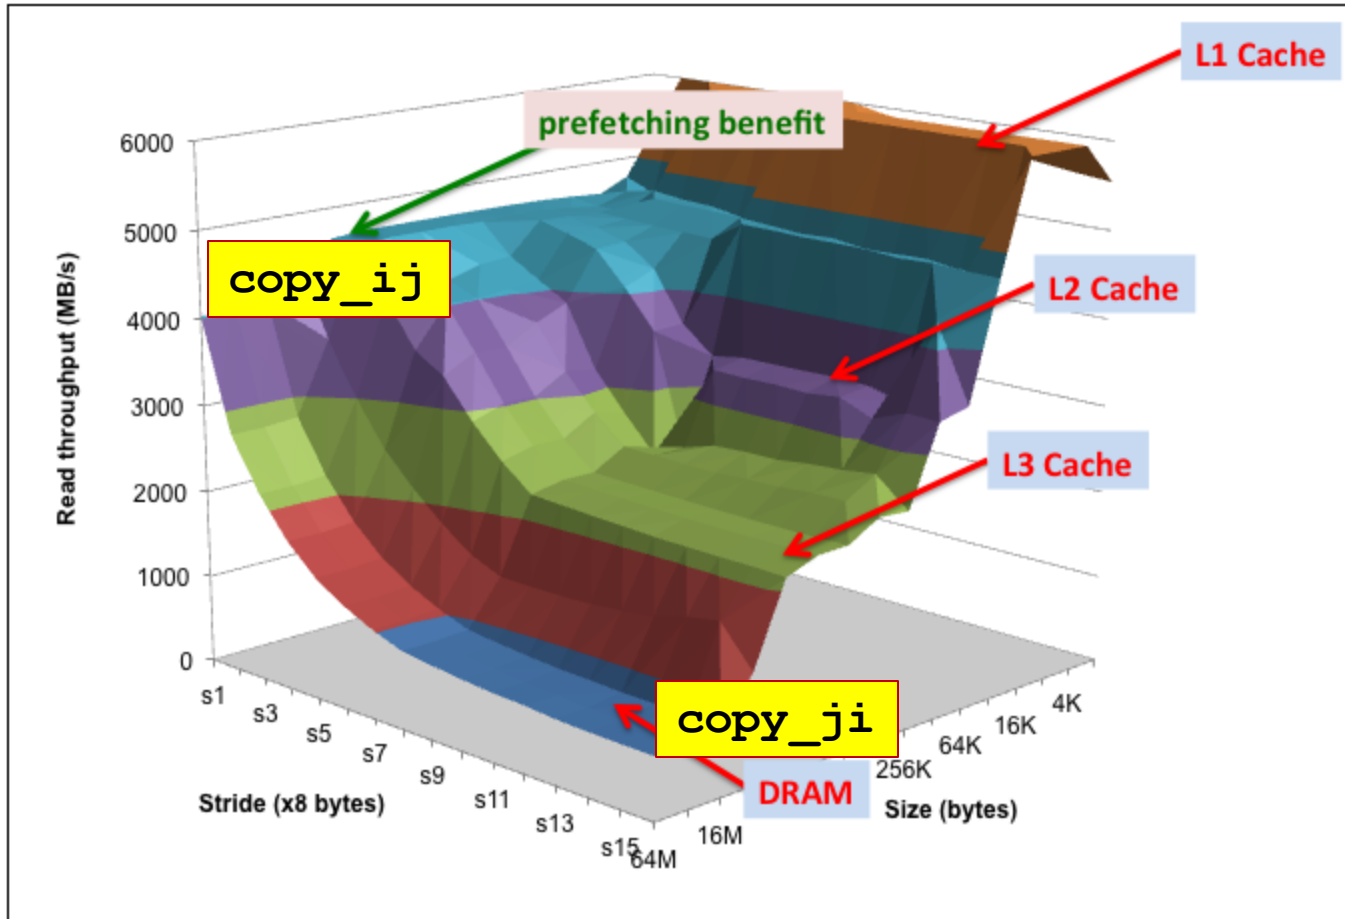
5.2 ms

```
void copy_ji(int src[2048][2048],
             int dst[2048][2048]) {
    int i, j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

162 ms



# The memory mountain



# Today: Machine Programming – Basics

- **What is an ISA?**
- History of x86
- The assembler language in context
- Introduction to assembly: Registers, operands
- Instruction format
- Memory addressing modes

# Instruction Set Architecture (ISA)

- **Definition 1**

**Architecture (Instruction Set Architecture (ISA))** the parts of a processor design that one needs to understand to write assembly.

- **Definition 2**

**Microarchitecture** is the implementation of the architecture.

- Example ISAs: x86, MIPS, ia64, VAX, Alpha, ARM, etc.

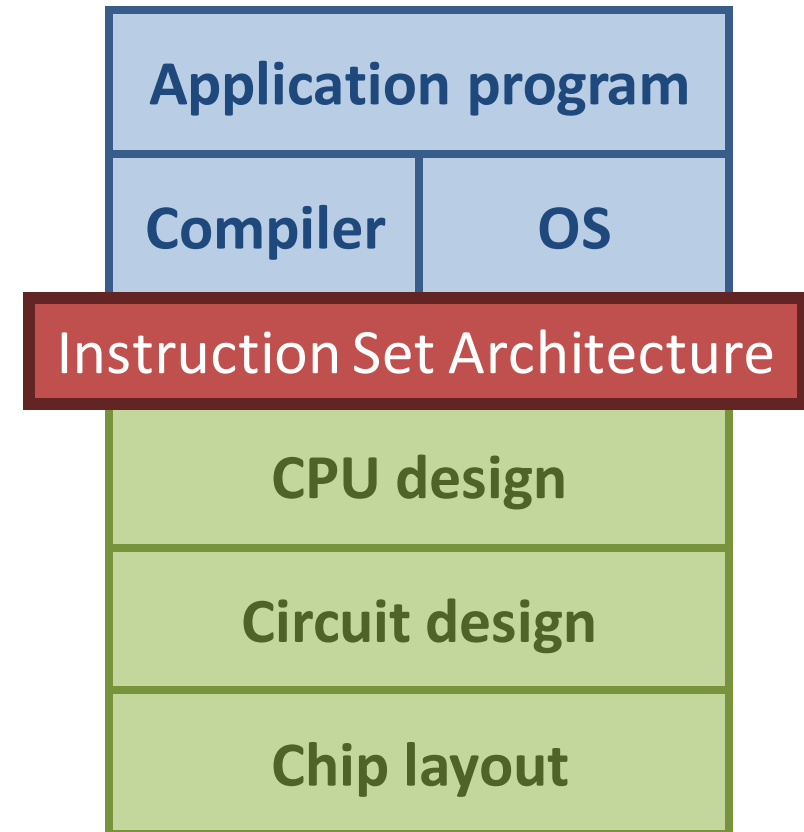
# Instruction Set Architecture – big picture

## ■ Assembly language view

- Processor state
  - Registers, memory, etc.
- Instructions
  - `addl`, `movq`, `leal`, etc.
  - How instructions are encoded as bytes

## ■ Layer of Abstraction

- Above: how to program a machine
  - Processor executes instructions in a sequence
- Below: what needs to be built
  - Use various tricks to make it run fast



# Today: Machine Programming – Basics

- What is an ISA?
- **Why x86?**
- The assembler language in context
- Introduction to assembly: Registers, operands
- Instruction format
- Memory addressing modes

# Intel x86 Processors

- **The x86 architecture dominates the laptop/desktop/server market**
- **Evolutionary design**
  - Backwards compatible up until 8086 (introduced in 1978)
  - Added more features as time goes on
- **Complex instruction set computer (CISC)**
  - Many different instructions with many different formats
  - Hard to match performance of RISC
  - But, Intel has done just that!


# Intel x86 Evolution: Key Milestones

<i>Name</i>	<i>Date</i>	<i>Transistors</i>	<i>MHz</i>
<b>8086</b>	<b>1978</b>	<b>29K</b>	<b>5-10</b>
<ul style="list-style-type: none"><li>■ First 16-bit processor. Basis for IBM PC &amp; DOS</li><li>■ 1 MB address space</li></ul>			
<b>386</b>	<b>1985</b>	<b>275K</b>	<b>16-33</b>
<ul style="list-style-type: none"><li>■ First 32-bit processor, referred to as IA32</li><li>■ Added “flat addressing”</li><li>■ Capable of running Unix</li></ul>			
<b>Pentium 4E</b>	<b>2004</b>	<b>125M</b>	<b>2800-3800</b>
<ul style="list-style-type: none"><li>■ First 64-bit Intel x86 processor, referred to as x86-64</li></ul>			

# Intel x86 Processors: Overview

## ■ Machine evolution, examples:

486	1989	1.9M
Pentium	1993	3.1M
PentiumPro	1995	6.5M
Pentium III	1999	8.2M
Pentium IV	2001	42M
Core 2 Duo	2006	291M
Xeon 7400	2008	1.9B
Xeon i7	2012	4.3B



## ■ Added features:

- transition from 32 bits to 64 bits
- more cores
- instructions to support multimedia operations, more efficient conditional ops



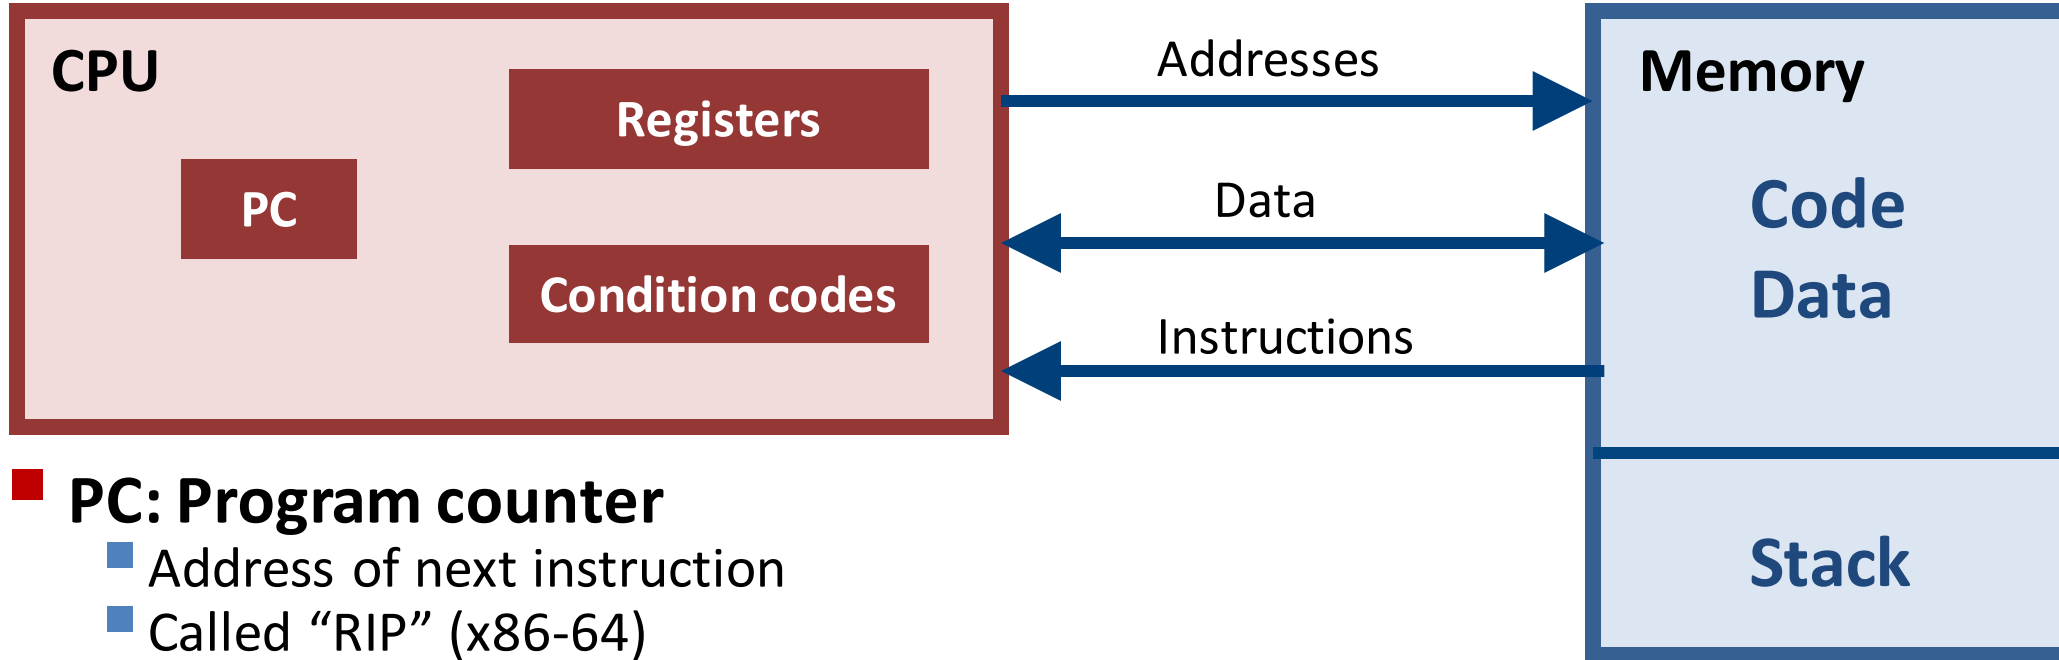
# A quick note on syntax

- There are two common ways to write x86 assembler code:
  - AT&T syntax
    - We are going to use it in this course
    - Common on Unix
  - Intel syntax
    - Generally used for Windows machines

# Today: Machine Programming – Basics

- What is an ISA?
- History of x86
- **The assembler language in context**
- Introduction to assembly: Registers, operands
- Instruction format
- Memory addressing modes

# Assembly/Machine code view



- **PC: Program counter**

- Address of next instruction
- Called “RIP” (x86-64)

- **Register file**

- Frequently used program data

- **Condition codes**

- Status information about recent arithmetic or logical operations
- Used for conditional branching

- **Memory**

- Byte addressable array
- code and user data
- stack to support procedures


# Compiling into Assembly

C code

```
int sum(int x, int y) {  
    int t = x + y;  
    return t;  
}
```

Generated x86 assembly

```
sum:  
    pushq    %rbp  
    movq     %rsp, %rbp  
    movl     %edi, -20(%rbp)  
    movl     %esi, -24(%rbp)  
    movl     -24(%rbp), %eax  
    movl     -20(%rbp), %edx  
    addl     %edx, %eax  
    movl     %eax, -4(%rbp)  
    movl     -4(%rbp), %eax  
    popq     %rbp  
    ret
```



- Running the command:  
**gcc -O -S code.c**
- Produces file: **code.s**

# Assembly characteristics:

## *Data Types and Instructions*

- **Data Types:**
  - Integers (1, 2, 4, 8 byte), Floating point (later in the course), no arrays/structs
- **Perform arithmetic function on register or memory data**
- **Transfer data between memory and register**
  - Load data from memory into a register
  - Store register data into memory
- **Transfer control**
  - Unconditional jumps to/from procedures
  - Conditional branches

# Machine instruction example

```
int t = x + y;
```

```
addl 8(%rbp), %eax
```

Similar to expression:

```
x += y
```

More precisely:

```
int eax;  
int *rbp;  
eax += rbp[2]
```

```
0x401046: 03 45 08
```

## ■ C code

- Add two integers

## ■ Assembly

- Add two 4-byte integers
  - “long” words in GCC parlance
- Operands:
  - **x**: register      **%eax**
  - **y**: memory      **M[%rbp+8]**
  - **t**: register      **%eax**
- Return function value in **%eax**

## ■ Object code

- 3-byte instruction
- stored at address **0x401046**

# Today: Machine Programming – Basics


- What is an ISA?
- History of x86
- The assembler language in context
- **Introduction to assembly: registers, data formats, operands**
- Instruction format
- Memory addressing modes

# What is a register?

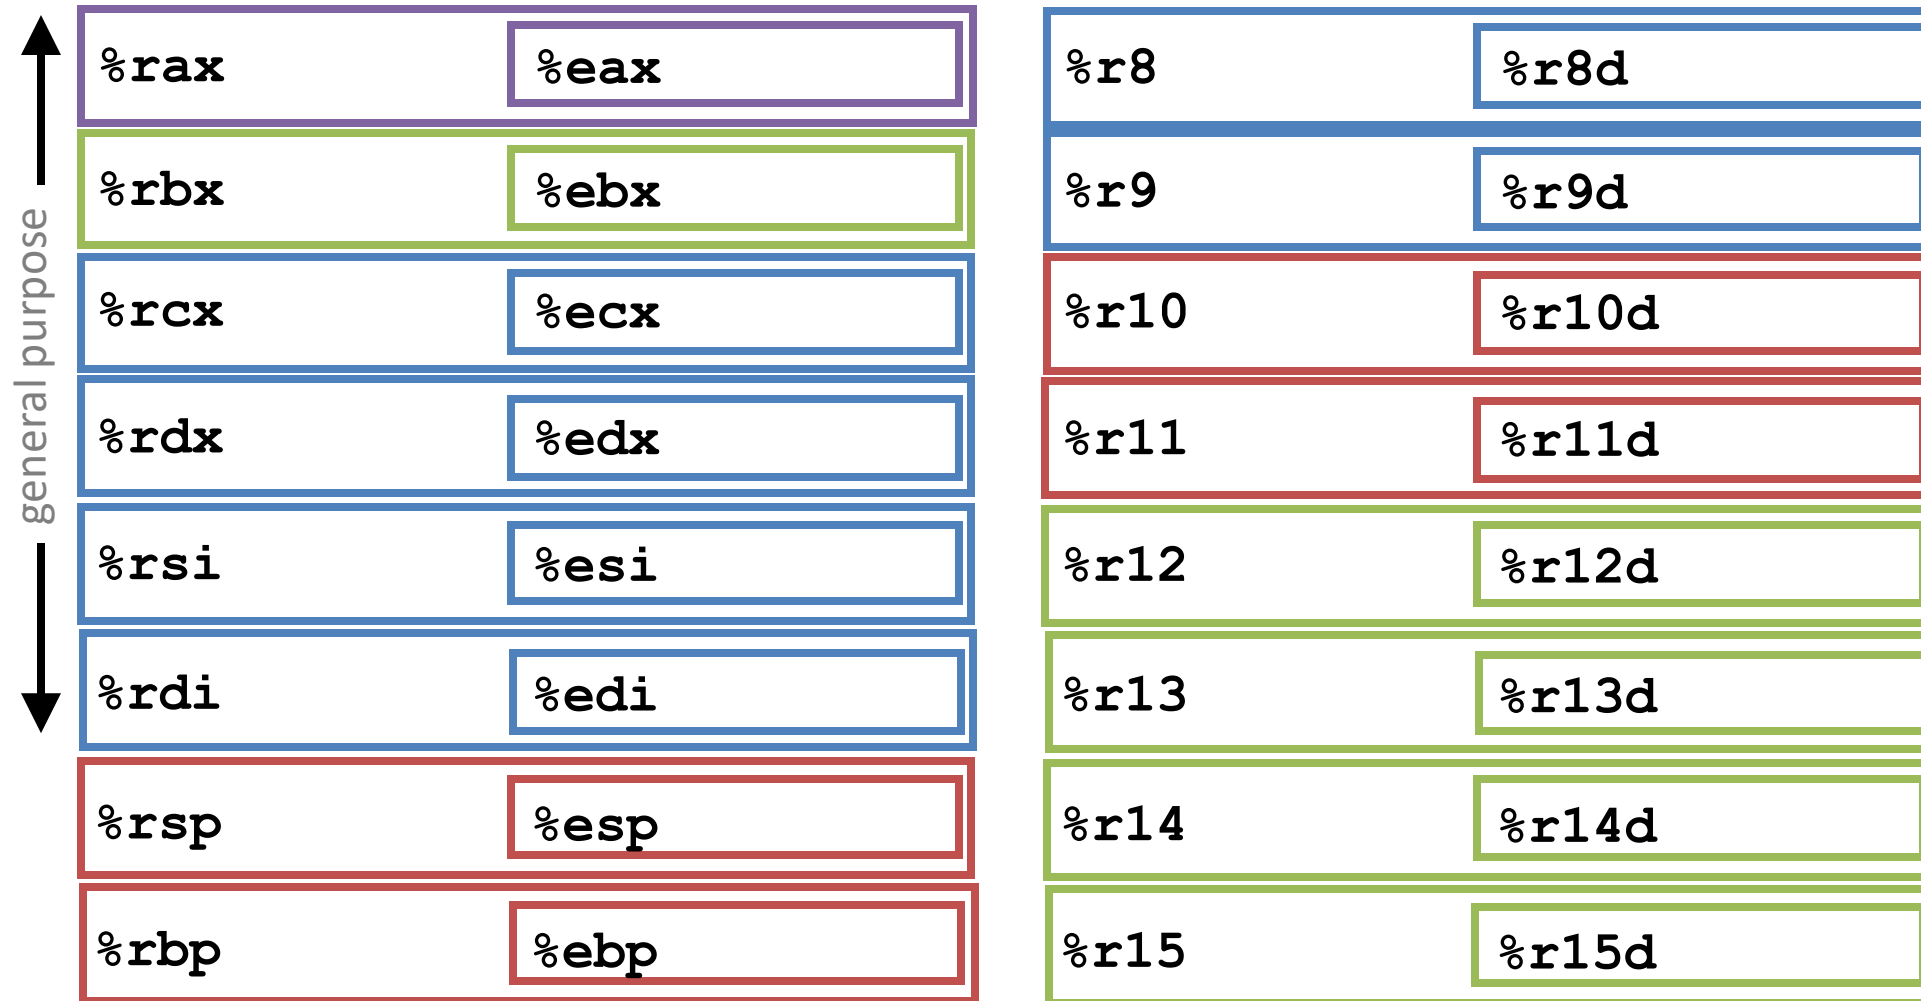
- A location in the CPU that stores a small amount of data, which can be accessed very quickly (once every clock cycle).
- Registers are at the heart of assembly programming.
  - They are a precious commodity in all architectures, but especially x86.



# History: IA32 Integer Registers



# x86-64 Integer Registers



# Data Formats

- Due to its origins as a 16-bit architecture, Intel uses ***word*** to refer to a 16-bit data type.
- 32-bit quantities are ***double words***, and 64-bit are ***quad words***.

C declaration	Intel data type	Assembly code suffix	Size (bytes)
<code>char</code>	Byte	<code>b</code>	1
<code>short</code>	Word	<code>w</code>	2
<code>int</code>	Double word	<code>l</code>	4
<code>long int</code>	Double word	<code>l</code>	4
<code>char*</code>	Double word	<code>l</code>	4
<code>float</code>	Single precision	<code>s</code>	4
<code>double</code>	Double precision	<code>q</code>	8
<code>long double</code>	Extended precision	<code>t</code>	10/12

# Instruction format

- Most Intel assembly *instructions* have at least one *operand*

label:	opcode	<i>source, destination</i>	;comments
label:	opcode	<i>operand</i>	;comments

- *label* is an optional user-defined identifier for the address of the instruction or data item which follows.
- The *operands* specify the source to reference in performing the operation and the destination location where to place the result.

# Operand Types

## 1. *Immediate* for constant values

Format:  $\$Imm$

Example:  $\$-536$ ,  $\$0x1F$

## 2. *Register* for the contents of one of the registers

Format: register  $E_a$ , referenced value is then  $R[E_a]$       Example:  $\%rax$ ,  $\%eax$

## 3. *Memory reference* to access a memory location based on a computed address

Format: memory address  $Addr$ , referenced value is then  $M[Addr]$

# Moving data

## ■ Moving Data

`movx Source, Destination`

## ■ Operand Types

- **Immediate:** constant integer data
  - Example: `$0x400`, `$-526`
  - Encoded with 1,2, or 4 bytes
- **Register:** one of 16 integer registers
  - Example: `%rax`, `%r13`
  - Note that `%rsp` is reserved for special use
- **Memory:** 8 consecutive bytes of memory at address given by register

<code>%rax</code>	<code>%eax</code>
<code>%rbx</code>	<code>%ebx</code>
<code>%rcx</code>	<code>%ecx</code>
<code>%rdx</code>	<code>%edx</code>
<code>%rsi</code>	<code>%esi</code>
<code>%rdi</code>	<code>%edi</code>
<code>%rsp</code>	<code>%esp</code>
<code>%rbp</code>	<code>%ebp</code>

<code>%r8</code>	<code>%r8d</code>
<code>%r9</code>	<code>%r9d</code>
<code>%r10</code>	<code>%r10d</code>
<code>%r11</code>	<code>%r11d</code>
<code>%r12</code>	<code>%r12d</code>
<code>%r13</code>	<code>%r13d</code>
<code>%r14</code>	<code>%r14d</code>
<code>%r15</code>	<code>%r15d</code>

# movl operand combinations

	Source	Dest	Source, Dest	C Analog
movl	Imm	Reg	movl \$0x4, %eax	temp = 0x4;
		Mem	movl \$-147, (%rax)	*p = -147;
	Reg	Reg	movl %eax, %edx	temp2 = temp1;
		Mem	movl %eax, (%rdx)	*p = temp;
	Mem	Reg	movl (%rax), %eax	temp = *p;

*Cannot do memory-to-memory transfer with a single instruction*

# Summary: Machine Programming – Basics

- What is an ISA?
- History of x86
- The assembler language in context
- Introduction to assembly: Registers, operands
- Instruction format
- **Memory addressing modes**



# Memory operands:

## Simple memory addressing modes

- The operand is the value at the specified address

- Normal addressing mode

- An Immediate (*Imm*) or a Register ( $R[E_a]$ ) specifies the memory address

$\text{Mem}[Imm]$  ,  $\text{Mem}[R[E_a]]$

Example: `movq (%rcx), %rax`

- Displacement

- An immediate (*Imm*) or a Register ( $R[E_a]$ ) specifies the start of a memory region
  - A constant displacement **D** specifies the offset

$\text{Mem}[Imm+D]$  ,  $\text{Mem}[R[E_a]+D]$

Example: `movq 8(%rbp), %rdx`

# Example of Simple Addressing Mode

C code

```
void swap(int* xp, int* yp) {  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

Compiled to assembly (with optimization)

```
swap:  
    movl (%rdi), %eax  
    movl (%rsi), %edx  
    movl %edx, (%rdi)  
    movl %eax, (%rsi)  
    ret
```

swap :

```
pushq %rbp  
movq %rsp, %rbp  
movq %rdi, -24(%rbp)  
movq %rsi, -32(%rbp)  
movq -24(%rbp), %rax  
movl (%rax), %eax  
movl %eax, -4(%rbp)  
movq -32(%rbp), %rax  
movl (%rax), %eax  
movl %eax, -8(%rbp)  
movq -24(%rbp), %rax  
movl -8(%rbp), %edx  
movl %edx, (%rax)  
movq -32(%rbp), %rax  
movl -4(%rbp), %edx  
movl %edx, (%rax)  
nop  
popq %rbp  
ret
```

# Example of Simple Addressing Mode


C code

```
void swap(int* xp, int* yp) {  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

Registers

%rdi	
%rsi	
%eax	
%edx	

Memory



Compiled to assembly (with optimization)

```
swap:  
    movl (%rdi), %eax  
    movl (%rsi), %edx  
    movl %edx, (%rdi)  
    movl %eax, (%rsi)  
    ret
```

# Example of Simple Addressing Mode

C code

```
void swap(int* xp, int* yp) {  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

Registers

%rdi	0x120
%rsi	0x110
%eax	
%edx	

Memory

Address

123	0x120
	0x118
456	0x110
	0x108
	0x100

Compiled to assembly (with optimization)

```
swap:  
    movl (%rdi), %eax  
    movl (%rsi), %edx  
    movl %edx, (%rdi)  
    movl %eax, (%rsi)  
    ret
```

# Example of Simple Addressing Mode

C code

```
void swap(int* xp, int* yp) {  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

Compiled to assembly (with optimization)

```
swap:  
    movl (%rdi), %eax  
    movl (%rsi), %edx  
    movl %edx, (%rdi)  
    movl %eax, (%rsi)  
    ret
```

Registers


%rdi	0x120
%rsi	0x110
%eax	123
%edx	

Memory

123
456

Address

0x120  
0x118  
0x110  
0x108  
0x100



# Example of Simple Addressing Mode

C code

```
void swap(int* xp, int* yp) {  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

Compiled to assembly (with optimization)

```
swap:  
    movl (%rdi), %eax  
    movl (%rsi), %edx  
    movl %edx, (%rdi)  
    movl %eax, (%rsi)  
    ret
```

Registers


%rdi	0x120
%rsi	0x110
%eax	123
%edx	456

Memory

123
456

Address

0x120  
0x118  
0x110  
0x108  
0x100



# Example of Simple Addressing Mode

C code

```
void swap(int* xp, int* yp) {  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

Compiled to assembly (with optimization)

```
swap:  
    movl (%rdi), %eax  
    movl (%rsi), %edx  
    movl %edx, (%rdi)  
    movl %eax, (%rsi)  
    ret
```


Registers

%rdi	0x120
%rsi	0x110
%eax	123
%edx	456

Memory

Address

456	0x120
	0x118
456	0x110
	0x108
	0x100



# Example of Simple Addressing Mode

C code

```
void swap(int* xp, int* yp) {  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

Compiled to assembly (with optimization)

```
swap:  
    movl (%rdi), %eax  
    movl (%rsi), %edx  
    movl %edx, (%rdi)  
    movl %eax, (%rsi)  
    ret
```

Registers

%rdi	0x120
%rsi	0x110
%eax	123
%edx	456

Memory

456
123

Address

0x120  
0x118  
0x110  
0x108  
0x100





# Memory operands:

## Complete memory addressing modes

### ■ Most General Form – D(Rb,Ri,S)

$$\text{Mem}[\text{R}[E_b] + s \cdot \text{R}[E_i] + D]$$

where

- $\text{R}[E_b]$  : is the base register: any of the 16 integer registers
- $\text{R}[E_i]$  : index register: any register, except for `%rsp`
- $D$  : constant displacement 1, 2, or 4 bytes (*why not 8?*)
- $S$  : scale: 1, 2, 4, or 8 (*why these numbers?*)

# Memory operands:

## Special cases memory addressing modes

### ■ Base register + index register

- Can be used to access **array elements** when **start of array is dynamically determined**

$\text{Mem}[\text{R}[E_b] + \text{R}[E_i]]$

where

$\text{R}[E_b]$  = the start of the array, and

$\text{R}[E_i]$  = byte index of array element

*Example assembly:*

```
mov (%bx,%di), %ax
```

Registers

%bx	0x8
%di	0x4
%ax	0x8

Memory

A[2]=8
A[1]
A[0]

Address

0x0c  
0x0a  
0x08  
...  
0x02  
0x00

# Memory operands:

## Special cases memory addressing modes

- **Relative based index = base register + index register + disp**
  - Can be used to access **arrays of objects**, **arrays within objects**, and **arrays on the stack**

$\text{Mem}[\text{R}[E_b] + \text{R}[E_i] + D]$

where

$\text{R}[E_b]$  = the start of the object

$D$  = array field within the object

$\text{R}[E_i]$  = byte offset of array element

*Example assembly:*

```
mov 0x8(%dx,%cx), %ax
```

Registers

%dx	0x2
%cx	0x4
%ax	0x5

Memory

A[2]=5
A[1]
A[0]

Address

0x0e  
0x0c  
0x0a  
...  
0x02  
0x00

# Memory operands:

## Special cases memory addressing modes

- **Base register + (scale \* index register) + displacement**
  - efficient access to **arrays elements within objects** and **on the stack** when the element size is 1, 2, 4 or 8 bytes

$$\text{Mem}[\text{R}[E_b] + s \cdot \text{R}[E_i] + D]$$

where

$\text{R}[E_b]$  = the start of the object

$D$  = array field within the object

$\text{R}[E_i]$  = “index” of array element

$s$  = element size

*Example assembly:*

```
mov 0x8(%dx,%cx,4), %ax
```

Registers

%dx	0x2
%cx	0x2
%ax	

Memory Address

A[2]	0x12
A[1]	0x10
A[1]	0x0e
A[0]	0x0c
A[0]	0x0a
...	...
	0x02
	0x00

# Summary: Machine Programming – Basics

- What is an ISA?
- History of x86
- The assembler language in context
- Introduction to assembly: Registers, operands
- Instruction format
- Memory addressing modes

# References

## ■ The slides are heavily influenced by:

- course “CS 15-213” at CMU (R. Bryant, D. O’Halloran, G. Kesden and M. Puschel)
- course “Computer Architecture and System Programming” at ETH Zurich (T. Roscoe)
- course “The Hardware/Software Interface” at UW (L. Ceze and G. Borriello)
- previous instalments of C113 (H. Wiklicky, A. Gopalan, R. Hayden and N. Dulay).

## ■ Additional information online:

- book “Computer Systems: A Programmer’s Perspective”, 3<sup>rd</sup> edition, Prentice Hall 2013
- online resources (related to the book, and else) <http://csapp.cs.cmu.edu>
- Intel manuals