

# 113: Architecture

Spring 2018

*Lecture:* X86 Data: Arrays and Structures

**Instructor:** Dr. Jana Giceva

# Today

- **Arrays**
  - One-dimensional
  - Multi-dimensional (nested)
  - Multi-level
- Structures

# Basic data types

## ■ Integral

- Stored and operated on in general (integer) registers
- Signed vs. unsigned depends on instructions used
- Example:
  - **byte** (size: 1 bytes, appendix: `b`, in C/Java: `char`)
  - **word** (size: 2 bytes, appendix: `w`, in C/Java: `short`)
  - **double word** (size: 4 bytes, appendix: `l`, in C/Java: `int`)
  - **quad word** (size: 8 bytes, appendix: `q`, in C/Java: `long int` (x86-64))

## ■ Floating point

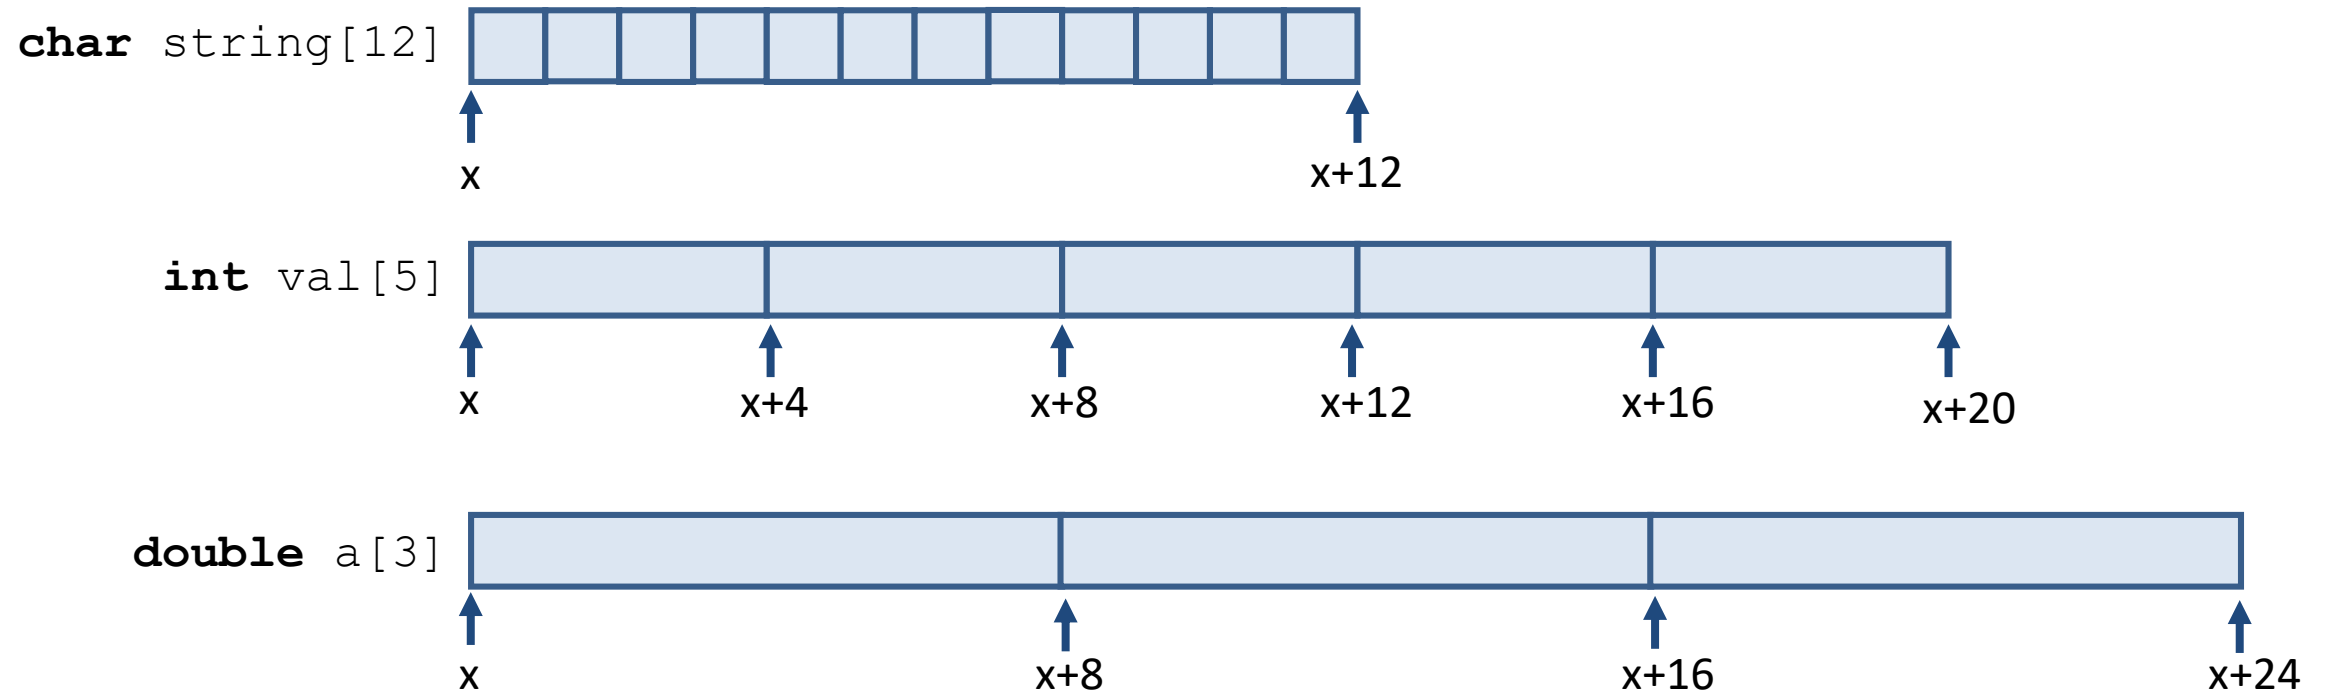
- Stored and operated on in floating point registers
  - Will cover them in more detail next week

# Array allocation


## ■ Basic Principle

$T$  **A**[ $L$ ];

- Array of data type  $T$  and length  $L$
- Contiguously allocated region of  $L * \text{sizeof}(T)$  bytes in memory



# Array Accessing Example



```
int get_digit(int val[], int dig)
{
    return val[dig];
}
```

```
get_digit:
    movslq    %esi, %rsi
    movl      (%rdi,%rsi,4), %eax
    ret
```

- Register `%rdi` contains starting address of array
- Register `%rsi` contains array index
- Desired digit at  $4 * \%rsi + \%rdi$
- Use memory reference  $(\%rdi, \%rsi, 4)$ .

# Array Loop Example


```
void digit_inc(int val[], int len){
    int i;
    for (i = 0; i < len; i++)
        val[i]++;
    return;
}
```

```
void digit_inc(int val[], int len){
    if (!len){
        int *vend = val + len-1;
        do {
            *ptr = *ptr + 1;
            ptr++;
        } while (ptr <= vend);
    }
    return;
}
```

```
digit_inc:
    testl    %esi, %esi
    jle      .L1
    movq     %rdi, %rax
    leal     -1(%rsi), %edx
    leaq     4(%rdi,%rdx,4), %rdx
.L3:
    addl     $1(%rax)
    addq     $4,%rax
    cmpq     %rdx, %rax
    jne      .L3
.L1:
    rep ret
```

# Nested Arrays

```
typedef int zip_dig[5];  
zip_dig pgh[4] =  
{ {1, 5, 2, 0, 6},  
  {1, 5, 2, 1, 3},  
  {1, 5, 2, 1, 7},  
  {1, 5, 2, 2, 1} };
```



- “`zip_dig pgh[4]`” equivalent to “`int pgh[4][5]`”
- Variable `pgh`: array of 4 elements, allocated contiguously
- Each element is an array of 5 `int`’s, allocated contiguously
- “**Row-Major**” ordering of all elements guaranteed



# Multidimensional (nested) Arrays

- Declaration
  - $T \ A[R][C];$
  - 2D array of data type  $T$
  - $R$  rows,  $C$  columns
  - Type  $T$  elements require  $K$  bytes

- Array size
  - $R * C * K$  bytes

- Arrangement
  - Row-Major ordering

`int A[R][C];`






# Nested Array Row Access

- Row vectors
  - $A[i]$  is array of  $C$  elements
  - Each element of type  $T$  requires  $K$  bytes
  - Starting address  $A + i * (C * K)$

```
int A[R][C];
```



# Nested Array Row Access Mode

```
int *get_pgh_zip(int index)
{
    return pgh[index];
}
```

```
zip_dig pgh[4] =
{{1, 5, 2, 0, 6},
 {1, 5, 2, 1, 3},
 {1, 5, 2, 1, 7},
 {1, 5, 2, 2, 1}};
```

```
leaq    (%rdi,%rdi,4), %rax    # 5 * index
leaq    pgh(,%rax,4), %rax     # pgh + (20 * index)
```

## ■ Row vector:

- `pgh[index]` is array of 5 `int`'s,
- starting address `pgh+20*index`


## ■ x86\_64 code:

- Computes and returns address
- Compute as `pgh+4*(index+4*index)`

# Nested Array Element Access

- Array Elements:
  - $A[i][j]$  is element of type  $T$ , which requires  $K$  bytes
  - $$\begin{aligned}\text{Address} &= A + i * (C * K) + j * K \\ &= A + (i * C + j) * K\end{aligned}$$

```
int A[R][C];
```



# Nested Array Element Access Code

```
int get_pgh_digit(int index, int digit ){
    return pgh[index][digit];
}
```

```
leaq    (%rdi, %rdi, 4), %rax    # 5 * index
addq    %rax, %rsi              # digit + 5 * index
movl    pgh(, %rsi, 4), %eax    # *(pgh + 4*(digit + 5*index))
```

## ■ Array Elements:

- `pgh[index][digit]` is `int` and `sizeof(int)=4`
- Address: `pgh + 5*4*index + 4*digit`

## ■ Assembly Code:

- Computes address as: `pgh + ((index+4*index) + digit)*4`
- `movl` performs memory reference


# Multi-Level Array Example

## Multi-Level Array Declaration(s):

```
int cmu[5] = {1, 5, 2, 1, 3};  
int mit[5] = {0, 2, 1, 3, 9};  
int ucb[5] = {9, 4, 7, 2, 0};
```

```
int* univ[3] = {mit, cmu, ucb};
```


- Variable `univ` denotes array of 3 elements
- Each element is a pointer: 8 bytes
- Each pointer points to array of `int`'s



Note: this is how Java represents multi-dimensional arrays!

# Element Access in Multi-Level Array

```
int get_univ_digit
(int index, int dig ){
    return univ[index][dig];
}
```



```
salq    $2, %rsi          # rsi = 4*digit
addq    univ(,%rdi,8), %rsi # p = univ[index] + 4*digit
movl    (%rsi), %eax       # return *p
```


## ■ Computation (x86\_64)

- Element access `Mem[Mem[univ+8*index]+4*dig]`
- Must do **two memory reads**:
  - First get pointer to row array
  - Then access element within array
- But, allows inner arrays to be **different lengths** (although not in this example)

# Array Element Accesses

## Nested Array


```
int get_pgh_digit
(int index, int digit ) {
    return pgh[index][digit];
}
```



**Mem**[pgh+20\*index+4\*dig]

## Multi-level Array

```
int get_univ_digit
(int index, int digit ) {
    return univ[index][digit];
}
```



**Mem**[**Mem**[univ+8\*index]+4\*dig]

Access looks similar, but x86 element access mode is different

# Data Structures in Assembly

- Arrays
  - One-dimensional
  - Multi-dimensional (nested)
  - Multi-level
- **Structures**
  - Alignment



# Structs in C


- Way of defining compound data types
- A structured group of variables, possibly including other `structs`

```
typedef struct {  
    int lengthInSeconds;  
    int yearRecorded;  
} Song;  
  
Song song1;  
song1.lengthInSeconds = 213;  
song1.yearRecorded     = 1994;  
  
Song song2;  
song2.lengthInSeconds = 214;  
song2.yearRecorded     = 1988;
```

- Given a struct instance, access member using the `.` operator:
- Given a *pointer* to a struct
- In assembly: pointer holds address of the first byte. Access elements with offsets.

# Structure Representation


```
typedef rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```



- Characteristics:
  - Contiguously-allocated region of memory
  - Refer to members within structure by names
  - Members may be of different types

# Structure Representation


```
typedef rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```



- **Structure represented as block of memory:**
  - Big enough to hold all of the fields
- **Fields ordered according to declaration order**
  - Even if another ordering would be more compact
- **Compiler determines overall size + positions of fields**
  - Machine-level programs has no understanding of the structures in the source code

# Accessing a Structure Member

```
typedef rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```




- Compiler knows the **offset** of each member within a struct
- Compute as
  - $* (r + \text{offset})$
  - Referring to absolute offset, so no pointer arithmetic

```
long get_i(struct rec *r)  
{  
    return r->i;  
}
```

```
# r in %rdi, index in %rsi  
movq 16(%rdi), %rax  
ret
```

# Exercise: Pointer to Structure Member

```
typedef rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```



```
long get_i(struct rec *r){  
    return &(r->i);  
}
```



```
# r in %rdi  
leaq    16(%rdi), %rax  
ret
```


```
long get_i(struct rec *r){  
    return &(r->next);  
}
```



```
# r in %rdi  
leaq    24(%rdi), %rax  
ret
```

# Generating Pointer to Array Element

```
typedef rec {  
    int a[4];  
    long i;  
    struct rec *next;  
} *r;
```



## ■ Generating Pointer to Array Element:

- Offset of each structure member determined at compile time
- Compute as:  $r + 4 * \text{index}$

```
long find_addr_of_array_elem  
    (struct rec *r, long index) {  
    return &r->a[index];  
}
```

```
# r in %rdi, index in %rsi  
leaq    (%rdi,%rsi,4), %rax  
ret
```

# Memory Alignment in x86-64

- For good memory system performance, Intel recommends data to be aligned
  - However, the x86-64 hardware will work correctly regardless of alignment of data
- *Aligned* means that any primitive object of  $K$  bytes must have an address that is multiple of  $K$

$K$	Type	Addresses
1	char	No restrictions
2	short	Lowest bit must be zero: $\dots 0_2$
4	int, float	Lowest 2 bits zero: $\dots 00_2$
8	long, double	Lowest 3 bits zero: $\dots 000_2$
16	long double	Lowest 4 bits zero: $\dots 0000_2$

# Alignment Principles

## ■ Aligned Data


- Primitive data type requires  $K$  bytes
- Address must be multiple of  $K$
- Required on some machines; advised on x86-64

## ■ Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
- Inefficient to load or store value that spans word boundaries

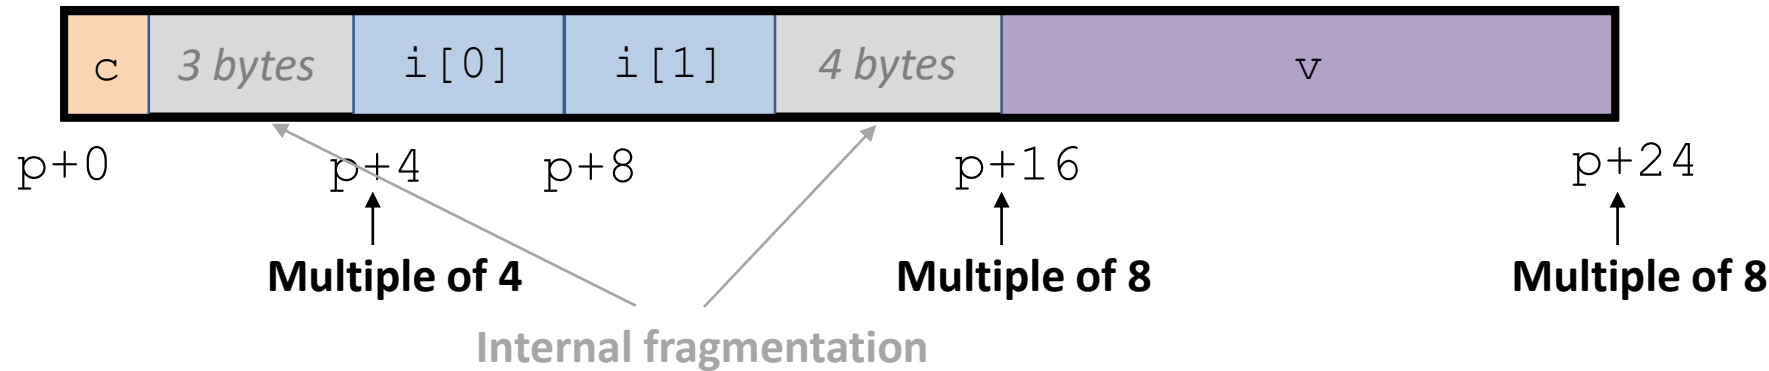


# Structures and Alignment



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

- **Aligned Data:**
  - Primitive data type requires  $K$  bytes
  - Address must be multiple of  $K$



# Satisfying Alignment with Structures

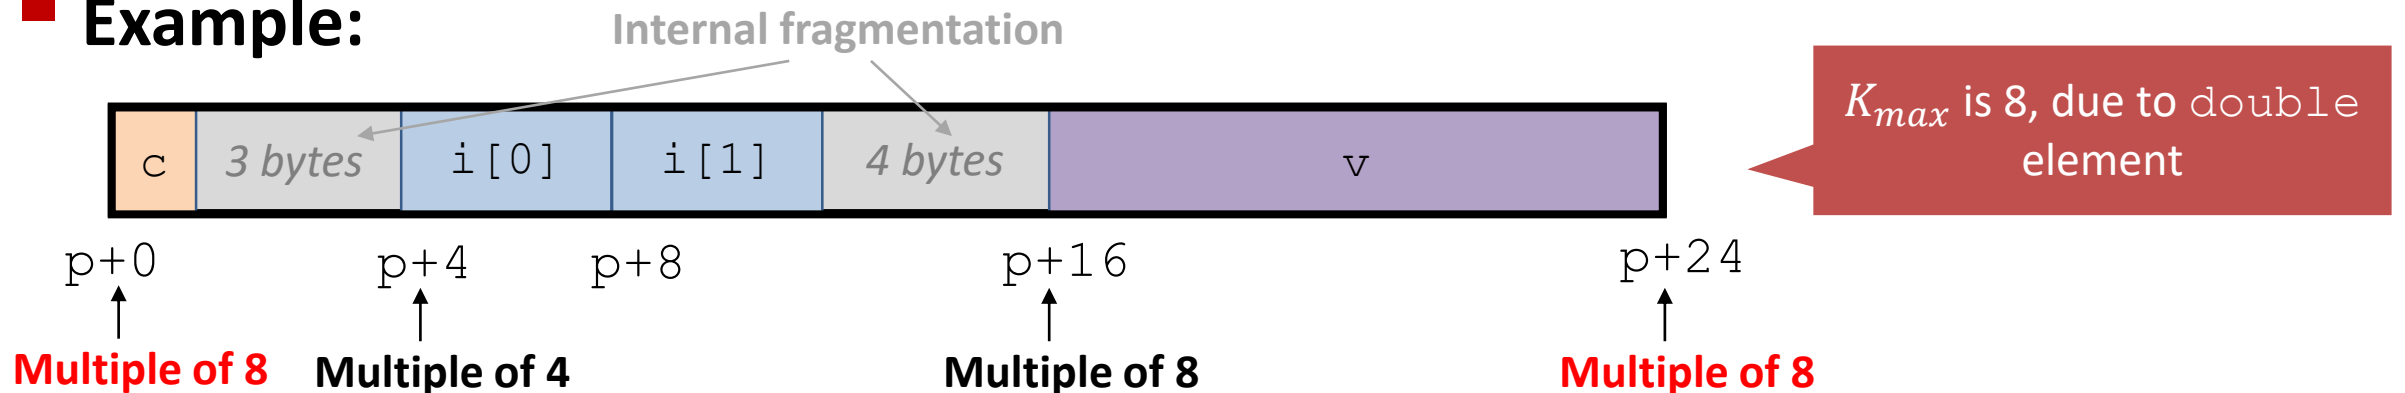
## ■ Within structure:

- Must satisfy each element's alignment requirement

## ■ Overall structure placement:

- Each structure has alignment requirement  $K_{max}$ 
  - $K_{max}$  = Largest alignment of any element
  - Counts array elements individually as elements
- Address of structure and structure length must be multiples of  $K_{max}$

## ■ Example:



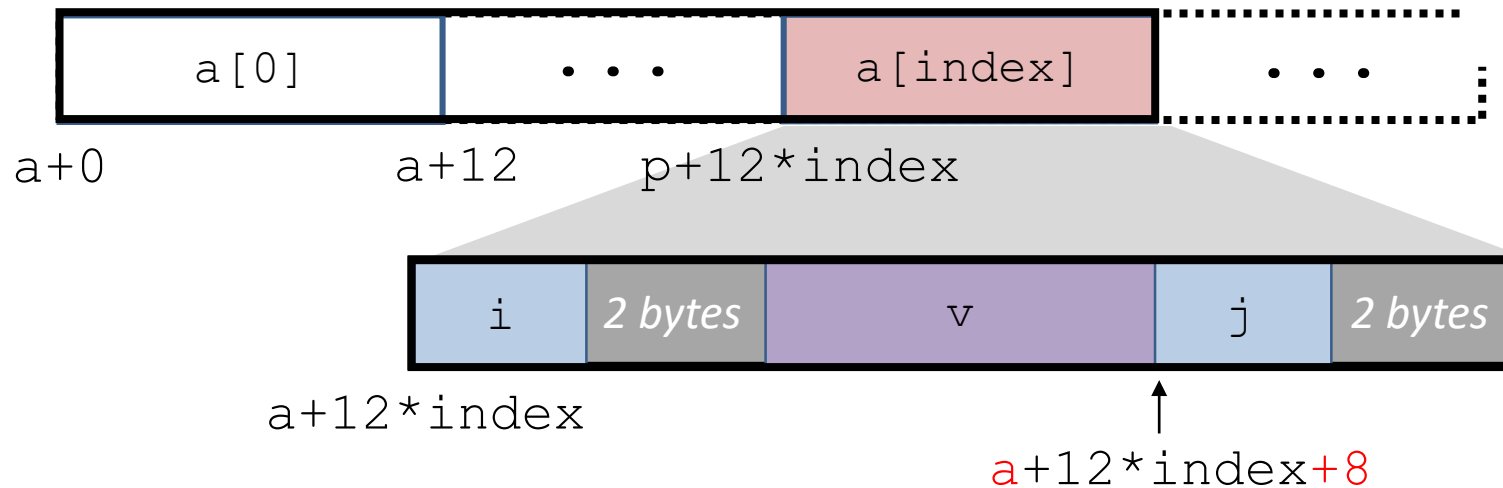
# Alignment of Structs

## ■ Compiler will do the following:

- Maintains declared ***ordering*** of fields in struct
- Each ***field*** must be aligned ***within*** the struct (*may insert padding*)
  - `offsetof` can be used to get actual field offset
- Overall struct must be ***aligned*** according to largest field
- Total struct ***size*** must be multiple of its alignment (*may insert padding*)
  - `sizeof` should be used to get true size of structs

# Accessing Array Elements within Structures

- Compute start of array element as:  $12 * \text{index}$ 
  - `sizeof(S3) = 12`, including alignment padding
- Element `j` is at offset 8 within structure
- Assembler gives offset `a+8`



```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```


```
short get_j(int index){  
    return a[index].j;  
};
```

```
# %rdi = index  
leaq (%rdi, %rdi, 2), %rax  
movzql a+8(, %rax, 4), %eax
```

# How the Programmer Can Save Space

- The compiler must respect the order elements are declared in
  - Sometimes the programmer can save space by declaring large data types first


```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



12 bytes



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```



8 bytes

# Summary

## ■ Arrays

- Contiguous allocations of memory
- Can usually be treated like a pointer to first element
- Nested arrays
  - all levels in one contiguous block of memory
- Multi-Level arrays
  - First level in one contiguous block of memory
  - Each element in the first level points to another “sub” array
  - Parts anywhere in memory

## ■ Structures

- Allocate bytes in order declared
- Padding in middle and at end to satisfy alignment