

NoSQL & NewSQL

Thomas Heinis

Imperial College
London

Some Data Management History...



The Database



What about...
Reliability?
Security?
Consistency?
Response time?
Scalability?



Relational Databases

- Structured, schema-based organization of data
- Data decomposed into tables, ex. customer data:

Customer		
Name	Date	City
Dustin	12.1.1971	London
Jack	8.19.1965	Leicester

City	
City	Country
London	UK
Leicester	UK

- SQL - General-purpose query language
- Join tables to retrieve all data: e.g., `SELECT * FROM Customer, City WHERE City.City = Customer.City`
- Focused on strong consistency (ACID)

Transactions – Data Integrity

Why Concurrent Access to Data must be managed?

John and Jane withdraw \$50 and \$100 from a common account...

John:

1. get balance
2. if balance > \$50
3. balance = balance - \$50
4. update balance

Jane:

1. get balance
2. if balance > \$100
3. balance = balance - \$100
4. update balance

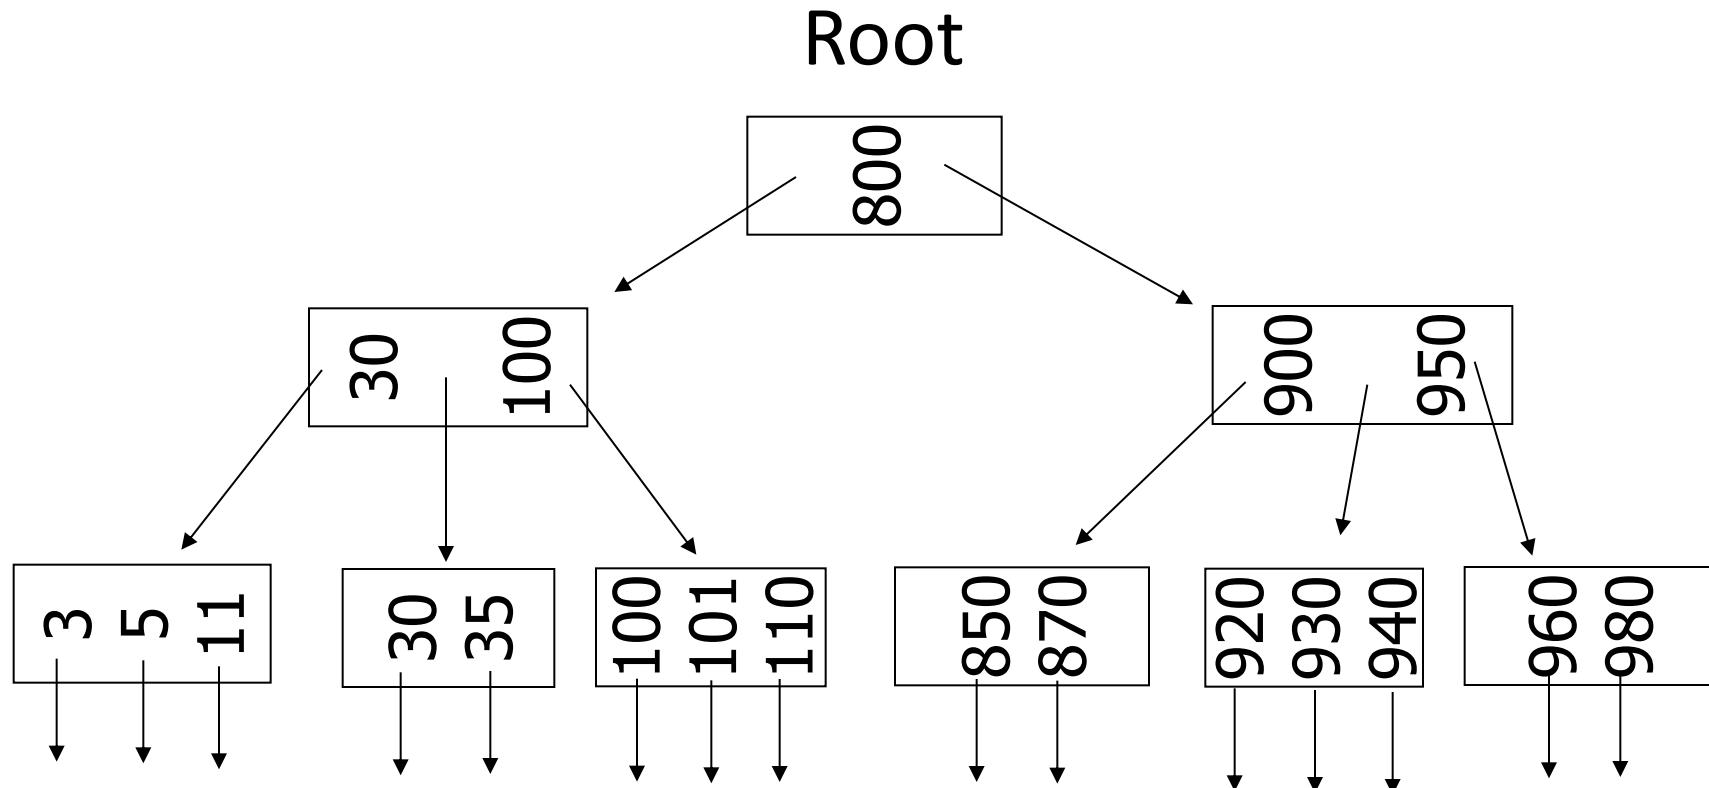
Initial balance \$300. Final balance=?

It depends...

Indexing – Efficient Data Retrieval

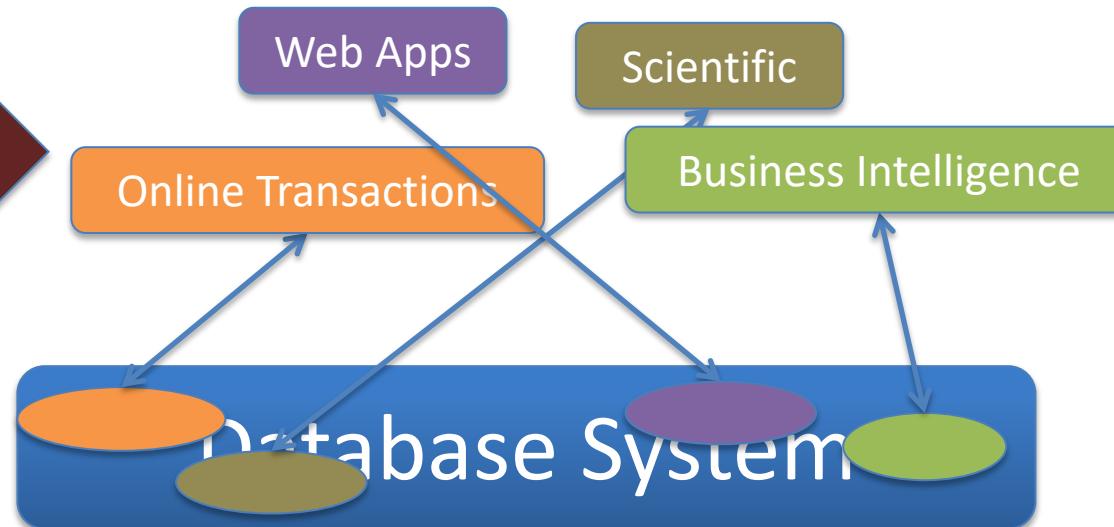
- How can we answer the query: “Find the account with the balance of 920.-” efficient?
- One approach is to scan the entire customer table, check every customer, return the one with balance = 920 ... very slow for large databases

Example Index (B-Tree)



Database Systems Today

New applications,
more requirements!



Rapidly changing hardware,
new performance hurdles!



Relational Databases: One Size Does Not Fit All...

New applications challenge relational databases:

1. Strong consistency (ACID) limits scalability
2. Schema evolution is challenging
3. Little optimization for novel hardware
4. Cumbersome language
5. Limited data types

NoSQL: Overview

- Main objective: implement distributed state
 - Different objects stored on different servers
 - Same object replicated on different servers
- Main idea: give up some of the ACID constraints to improve performance
- Simple interface:
 - Write (=Put): needs to write all replicas
 - Read (=Get): may get only one
- Eventual consistency ← Strong consistency

NoSQL

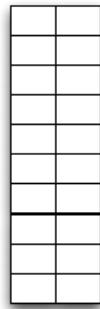
“Not Only SQL” or “Not Relational”.

Six key features:

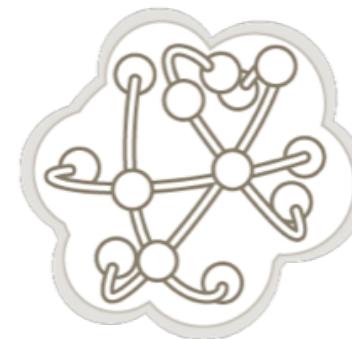
1. Scale horizontally “simple operations”
2. Replicate/distribute data over many servers
3. Simple call level interface (contrast w/ SQL)
4. Weaker concurrency model than ACID
5. Efficient use of distributed indexes and main memory
6. Flexible schema

Four NOSQL Categories

Key-Value



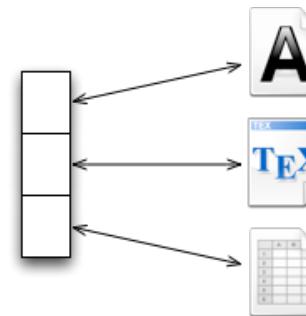
Graph DB



BigTable

				1					
					1				
		1				1			
							1		
			1					1	
									1
									1

Document



Data Model

- **Tuple** = row in a relational DB
- **Document** = nested values, extensible records
(think XML or JSON)
- **Extensible record** = families of attributes have a schema, but new attributes may be added
- **Object** = like in a programming language, but without methods

Key-value Stores

- Think “file system” more than “database”
- Consistent hashing (DHT)
- Only primary index: lookup by key
- No secondary indexes
- Transactions: single- or multi-update

Basic Idea: Key-Value Store

Table T:

key	value
k1	v1
k2	v2
k3	v3
k4	v4

keys are sorted



- API:
 - `lookup(key) → value`
 - `lookup(key range) → values`
 - `getNext → value`
 - `insert(key, value)`
 - `delete(key)`
- Each row has timestamp
- Single row actions atomic
No multi-key transactions
- No query language!

Document Stores

- A "document" = a pointerless object = e.g. JSON = nested or not = schema-less
- In addition to KV stores, may have secondary indexes
- SimpleDB, CouchDB, MongoDB, Terrastore
- Scalability:
 - Replication (e.g. SimpleDB, CouncillDB – means entire db is replicated),
 - Sharding (MongoDB);
 - Both

Document Store (MongoDB)

```
> db.user.insert({  
    first: "John",  
    last : "Doe",  
    age: 39  
})
```

```
> db.user.find ({"first" : "John"})  
{  
    "_id" : ObjectId("51..."),  
    "first" : "John",  
    "last" : "Doe",  
    "age" : 39  
}
```

```
> db.user.update(  
    {"_id" : ObjectId("51...")},  
    {  
        $set: {  
            age: 40,  
            salary: 7000}  
    }  
)
```

```
> db.user.remove({  
    "first": /^J/  
})
```

Column Family

- Most Based on **BigTable**: Google's Distributed Storage System for Structured Data
- Data Model:
 - A big table, with column families
 - Map Reduce for querying/processing
- Examples:
 - HBase, HyperTable, Cassandra

Column Family: Pros and Cons

- Pros:
 - Supports Semi-Structured Data
 - Naturally Indexed (columns)
 - Scalable
- Cons
 - Poor for interconnected data

Graph Databases

- Data Model:
 - Nodes and Relationships
- Examples:
 - Neo4j, OrientDB, InfiniteGraph, AllegroGraph

Graph Databases: Pros and Cons

- Pros:
 - Powerful data model, as general as RDBMS
 - Connected data locally indexed
 - Easy to query
- Cons
 - Sharding (lots of people working on this)
 - Scales UP reasonably well
 - Requires rewiring your brain

What are graphs good for?

- Recommendations
- Business intelligence
- Social computing
- Geospatial
- Systems management
- Web of things
- Genealogy
- Time series data
- Product catalogue
- Web analytics
- Scientific computing (especially bioinformatics)
- Indexing your *slow* RDBMS
- And much more!

What is a Graph?

An abstract representation of a set of objects where some pairs are connected by links.



Object (Vertex, Node)



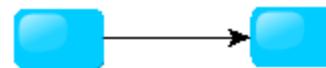
Link (Edge, Arc,
Relationship)

Different Kinds of Graphs

- Undirected Graph



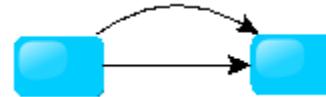
- Directed Graph



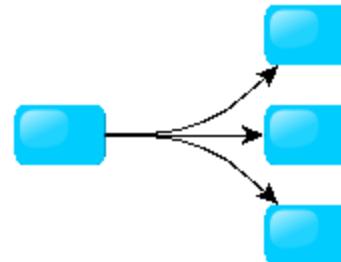
- Pseudo Graph



- Multi Graph



- Hyper Graph

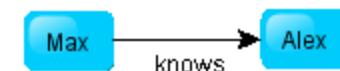


More Kinds of Graphs

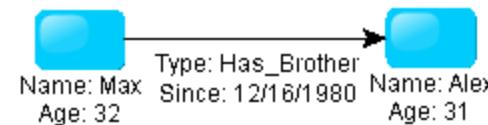
- Weighted Graph



- Labeled Graph



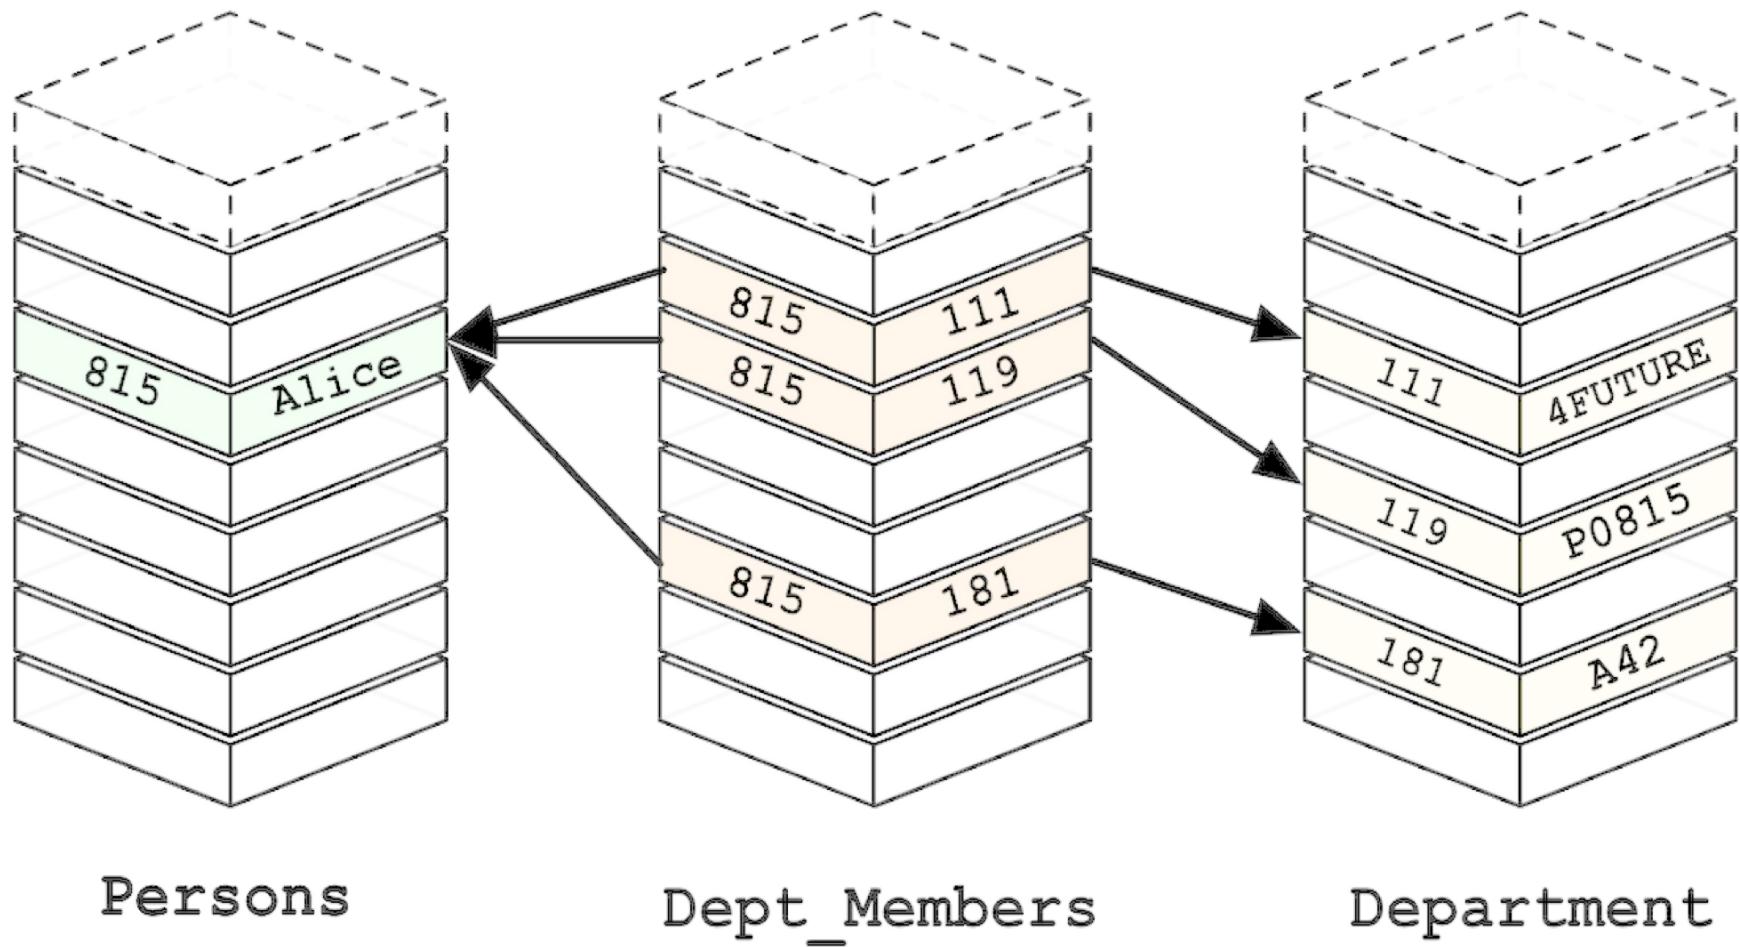
- Property Graph



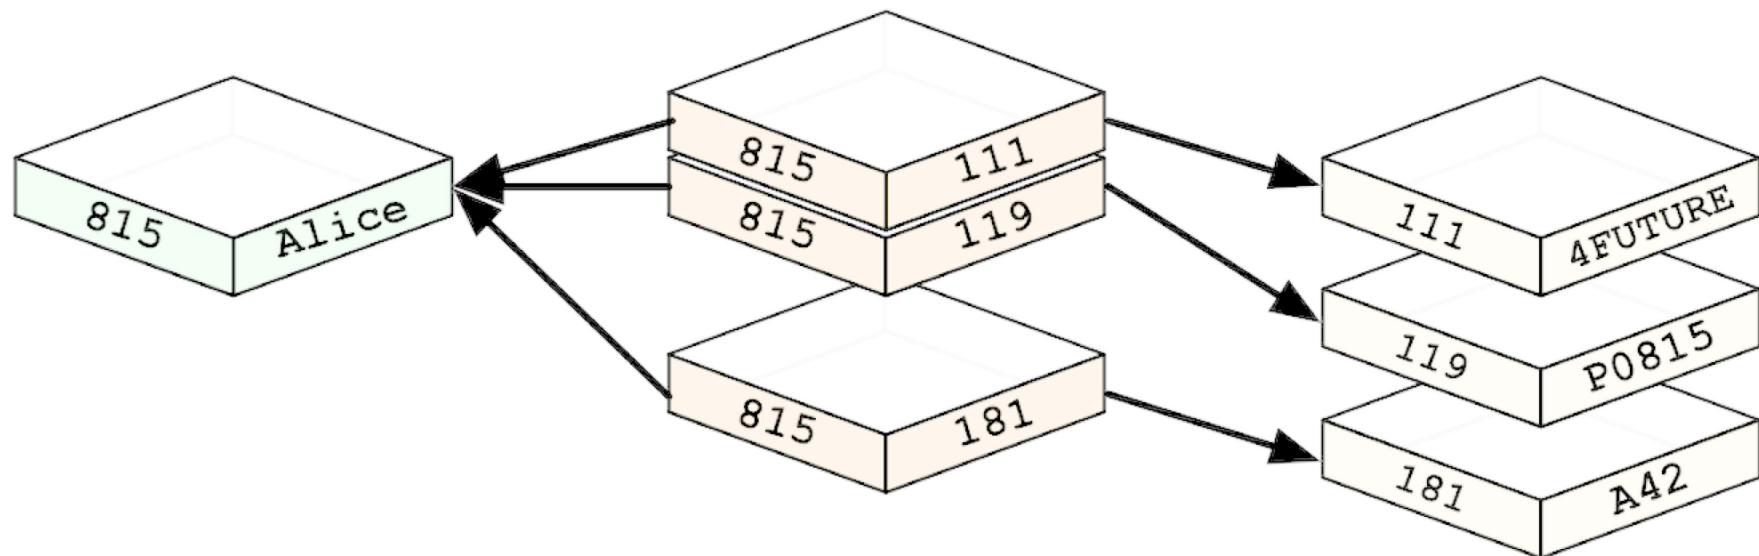
What is a Graph Database?

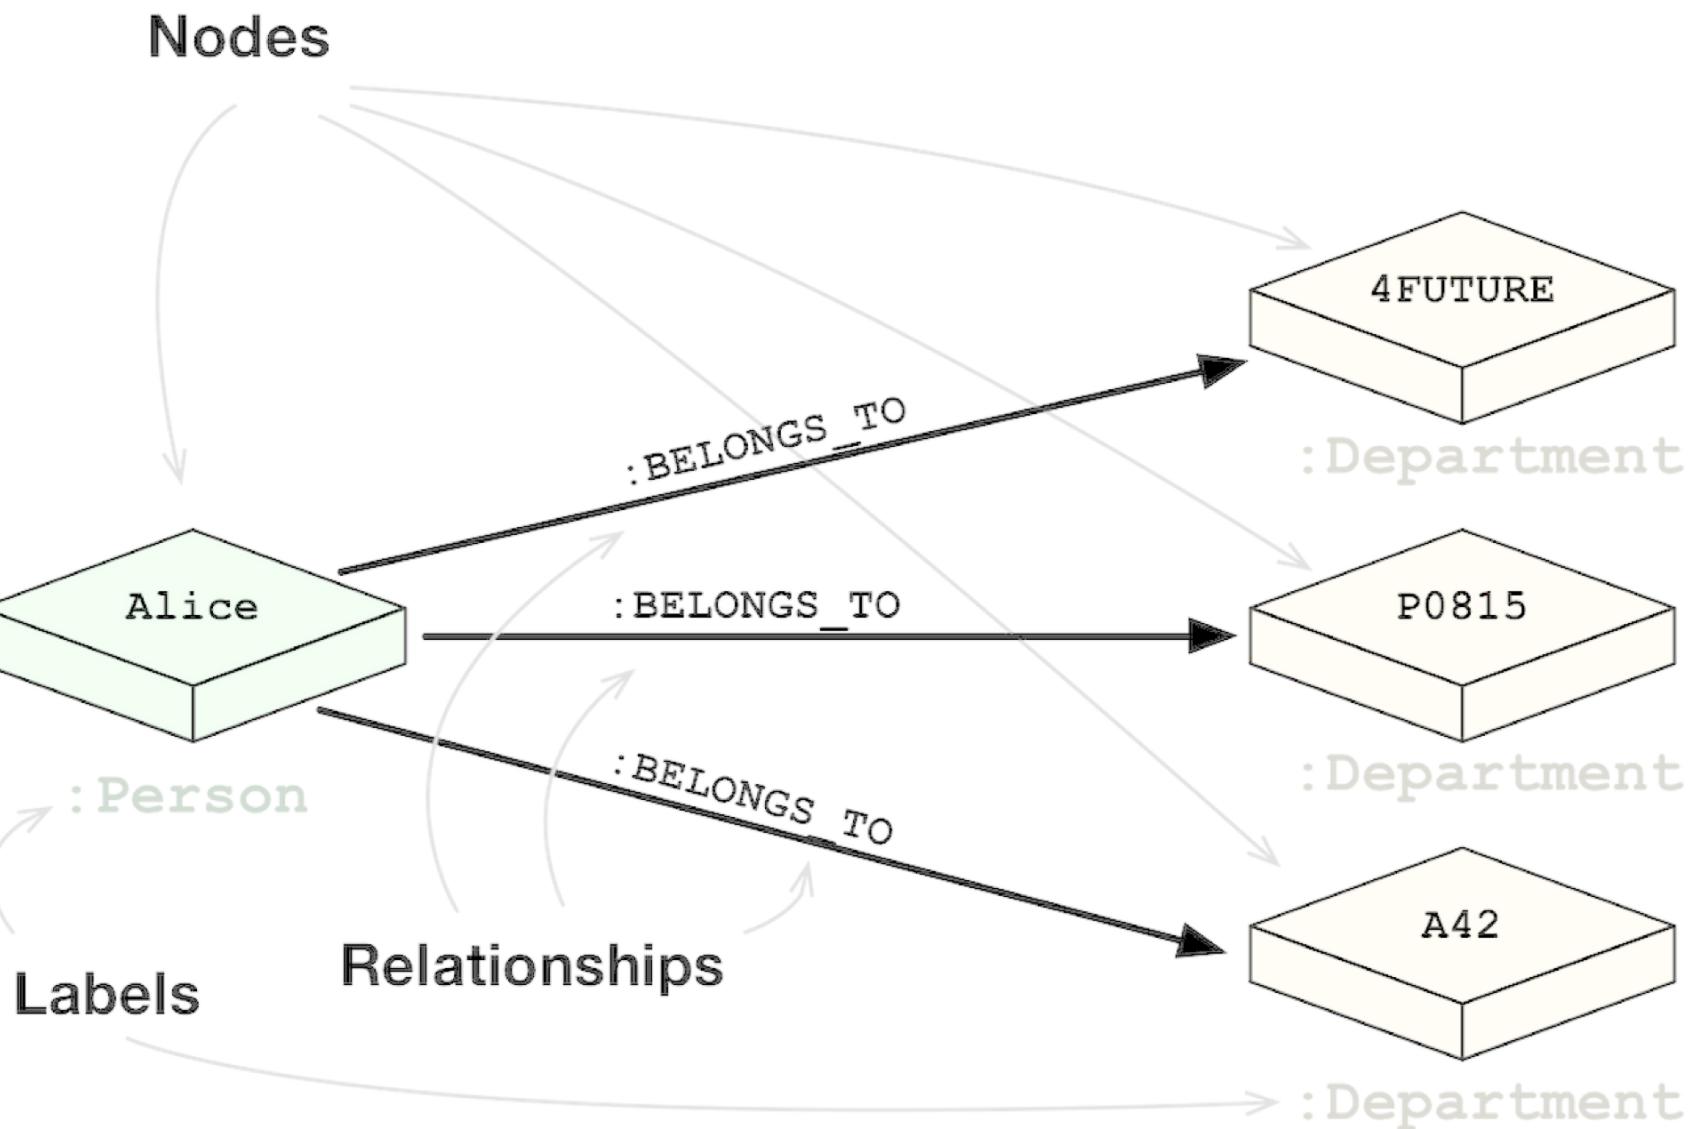
- A database with an explicit graph structure
- Each node knows its adjacent nodes
- As the number of nodes increases, the cost of a local step (or hop) remains the same
- Plus an Index for lookups

Relational Databases



Graph Databases

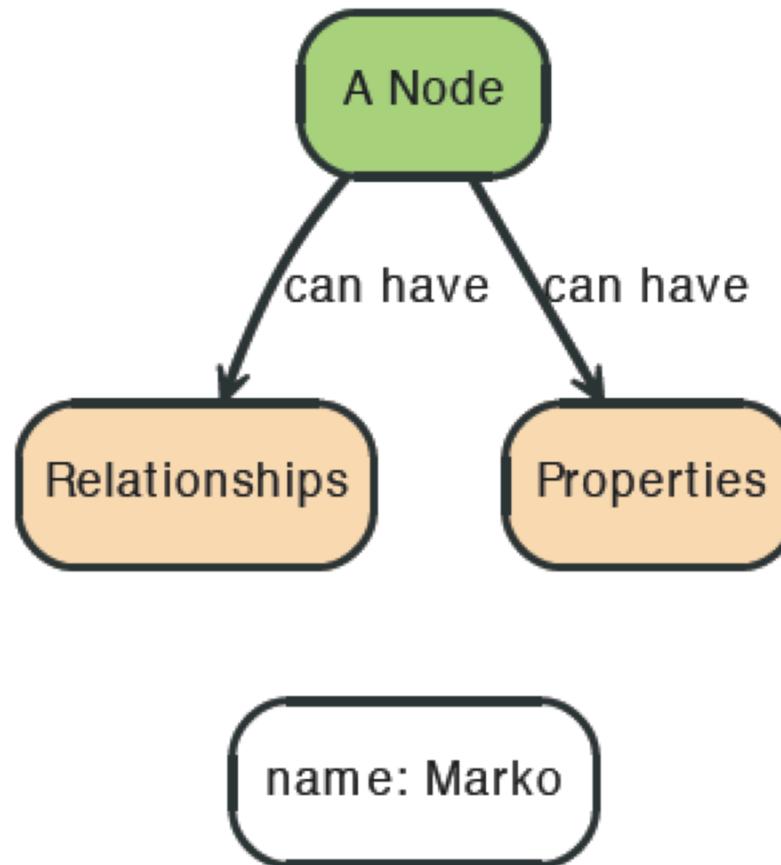




Neo4j Approach

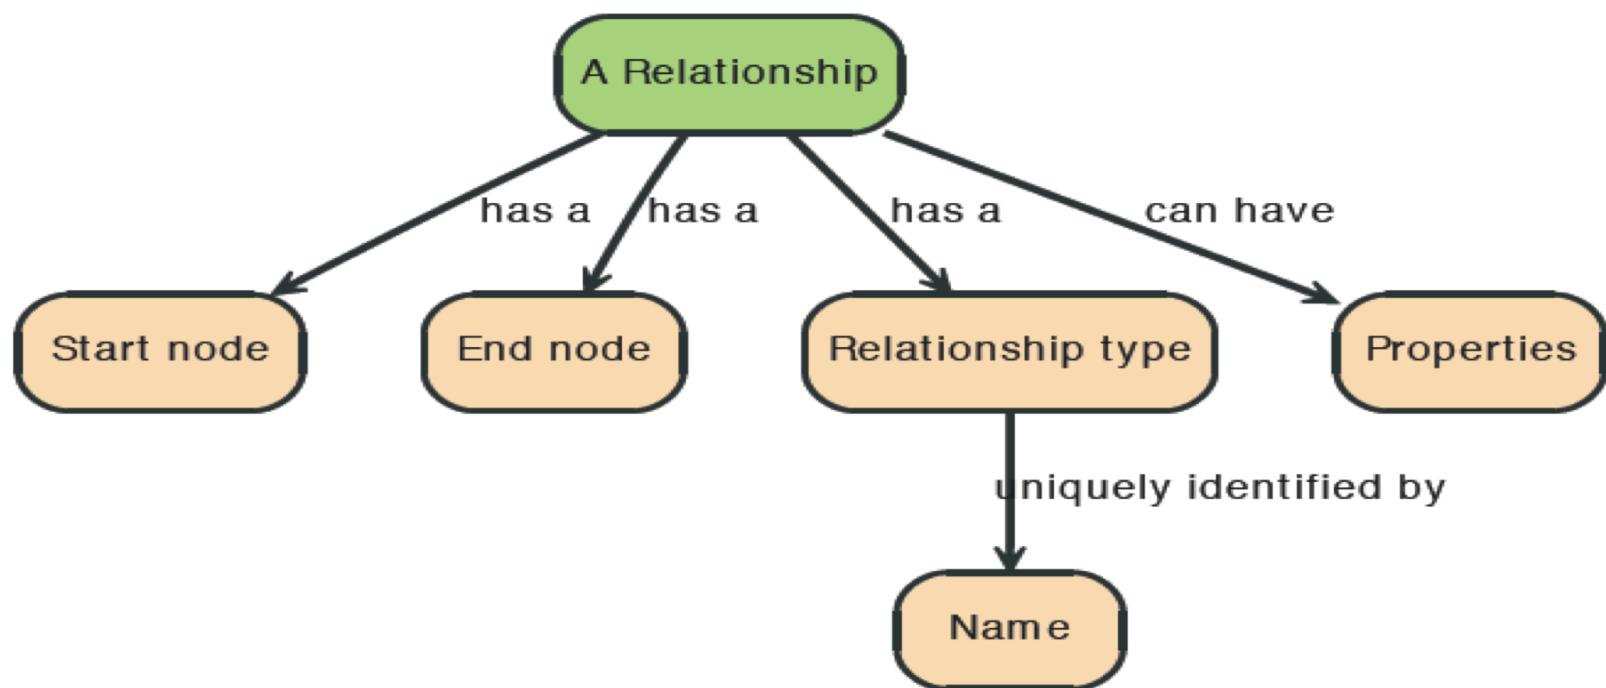
- Each entity table is represented by a label on nodes
- Each row in a entity table is a node
- Columns on those tables become node properties.
- Remove technical primary keys, keep business primary keys
- Add unique constraints for business primary keys, add indexes for frequent lookup attributes

Node in Neo4j

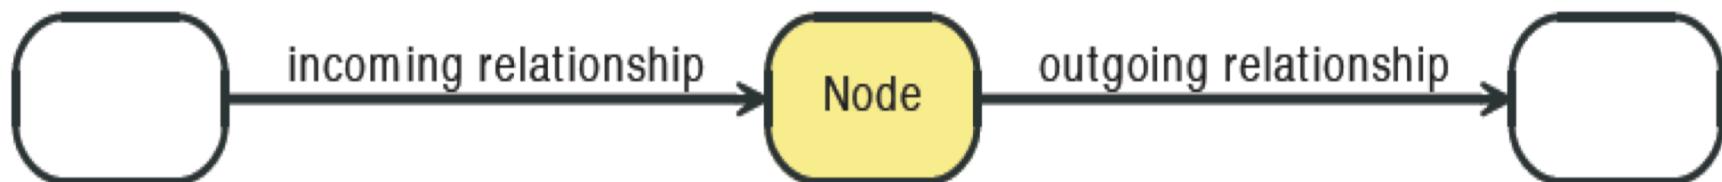
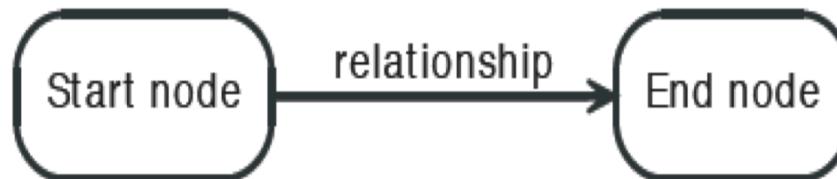


Relationships in Neo4j

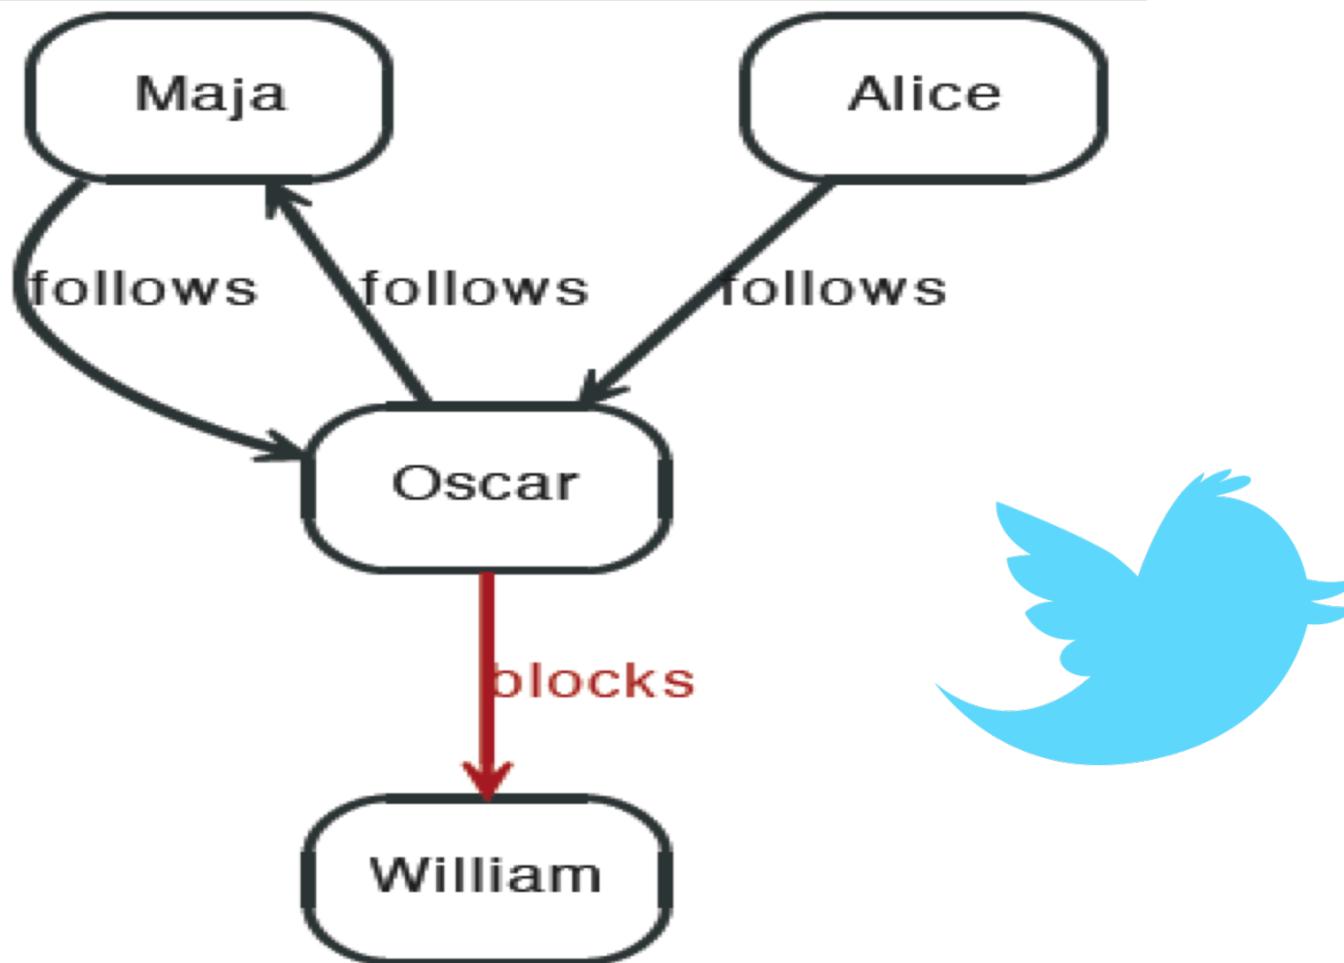
- Relationships between nodes are a key part of Neo4j.



Relationships in Neo4j



Twitter and relationships



(a:Person) – [:follows] -> (b:Person) – [:follows] -> (c:Person)

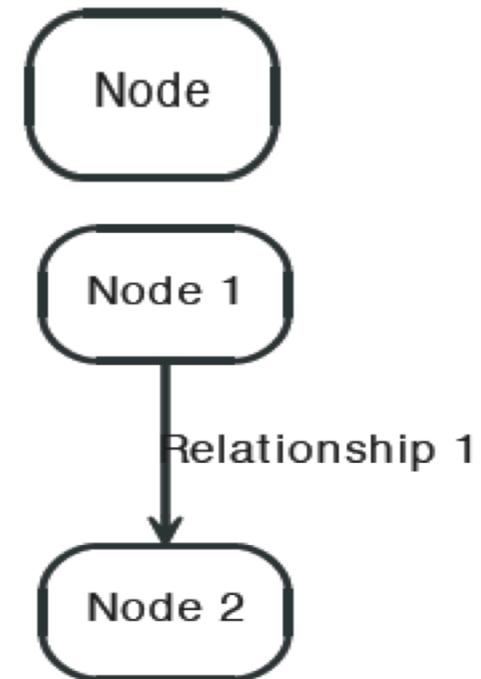
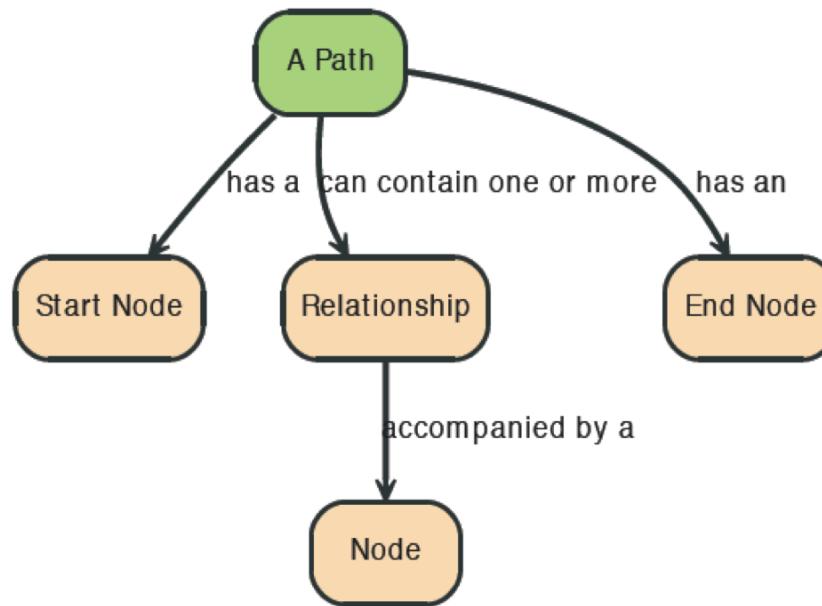
Properties

- Both nodes and relationships can have properties.
- Properties are key-value pairs where the key is a string.
- Property values can be either a primitive or an array of one primitive type.

For example String, int and int[] values are valid for properties.

Paths in Neo4j

A path is one or more nodes with connecting relationships, typically retrieved as a query or traversal result.



Scalable Relational Systems (NewSQL)

- Means RDBMS that offer sharding
- **Key difference:**
 - NoSQL difficult or impossible to perform large-scope operations and transactions (to ensure performance)
 - Scalable RDBMS do not **preclude** these operations, but users pay a price only when they need them
- MySQL Cluster, VoltDB, Clusterix, ScaleDB, Megastore (the new BigTable)
- Many more **NewSQL** systems becoming available...

Scalable Data Processing

- Parallel execution achieves greater efficiency
- But, parallel programming is hard
 - Parallelization
 - Fault Tolerance
 - Data Distribution
 - Load Balancing

MapReduce (Hadoop and others)

- “*MapReduce is a programming model and an associated implementation for processing and generating large data sets*”
- Programming model
 - Abstractions to express simple computations
- Library
 - Takes care of the gory stuff: Parallelization, Fault Tolerance, Data Distribution and Load Balancing

NoSQL Systems

- Key Value Stores

Key	Value
SW3 3TB	London
B18 4BJ	Birmingham



- Document Stores



- Scalable SQL Systems



- Data Processing Systems

NoSQL Challenges

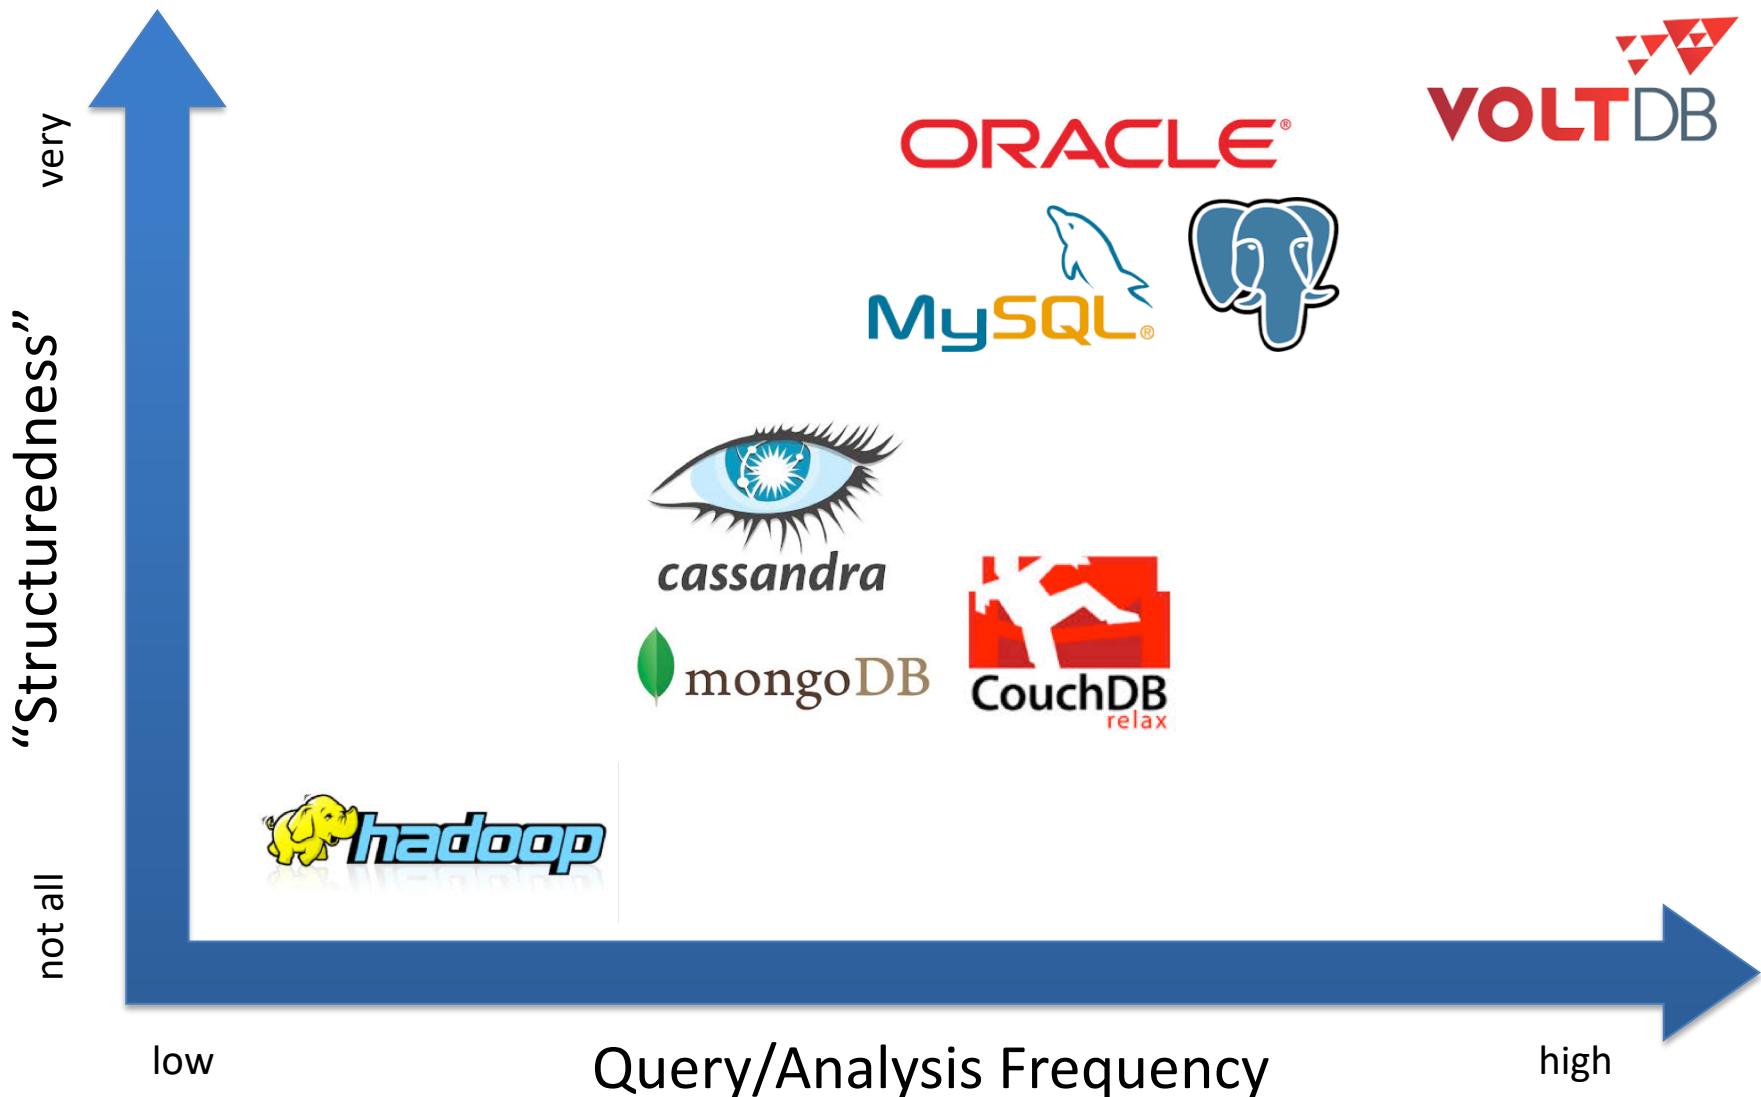
E.g., Flexible Schemas in MongoDB

```
db.inventory.insert(  
  {  
    category: "vacuum",  
    details: {  
      model: "14Q3",  
      manufacturer: "XYZ Company"  
    },  
    stock: [ { size: "S", qty: 25 }],  
    category: "clothing"  
  }  
)
```

```
db.inventory.insert(  
  {  
    category: "vacuum",  
    details: {  
      model: "14Q3",  
      manufacturer: "XYZ Company"  
    },  
    color: "blue"  
  })
```

What fields does db.user.find ({"category" : "vacuum"}) have?

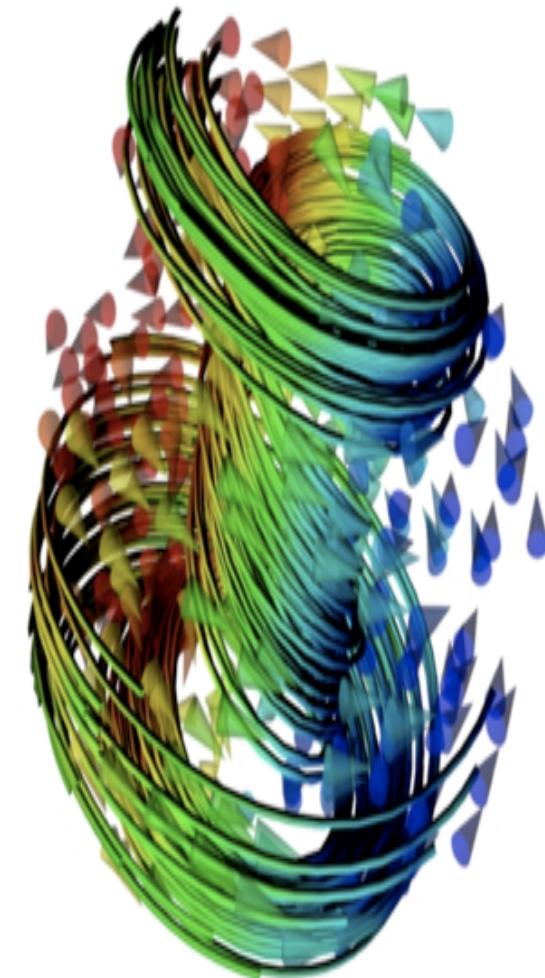
Data Management Landscape*



* according to me

What about Scientific Data?

- Cannot be represented as a table → different data models
- Do not need the abstraction of transaction: typically read-only
- Updates rare
- Massive in size → Indexing considerable overhead



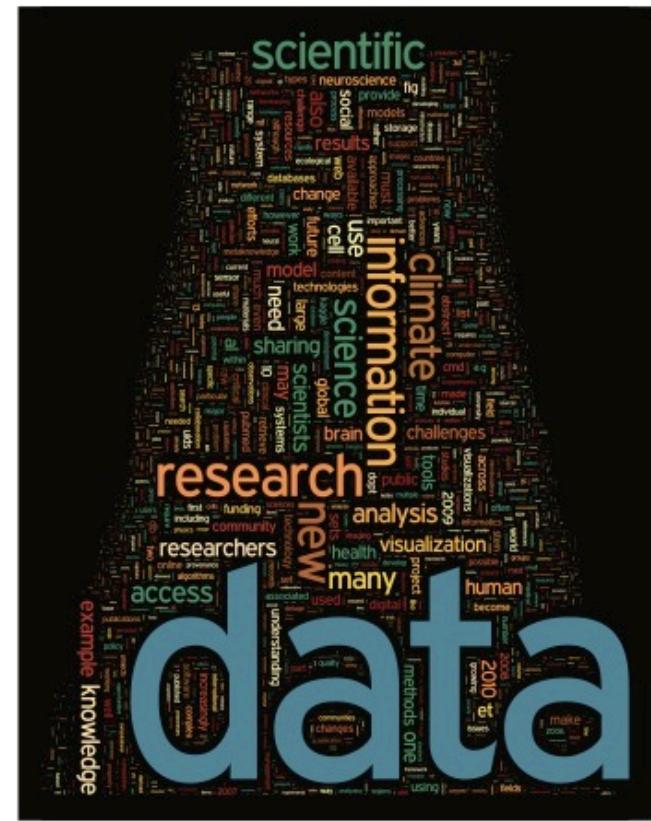
Many database optimizations irrelevant!

Data-driven Science

Past:

- theory
 - simulation
 - experiments

The “fourth paradigm” scientific breakthrough through computing on massive data



Scientific Data Management

Today:

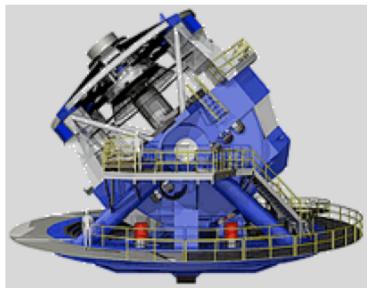
- legacy software
- in main memory of supercomputers
- databases too rigid to use

But as data grows, problem *changes*:

- difficult and slow
- some data discarded

Astronomy

~30,000 movies,



20 TB/day!



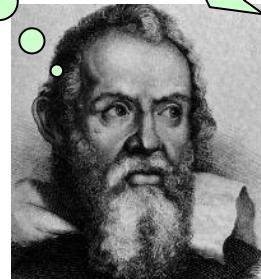
Table of Observations

	TYPE	CX	CY	E1	E2	E3	...

Large Synoptic Survey
Telescope (www.lsst.org)

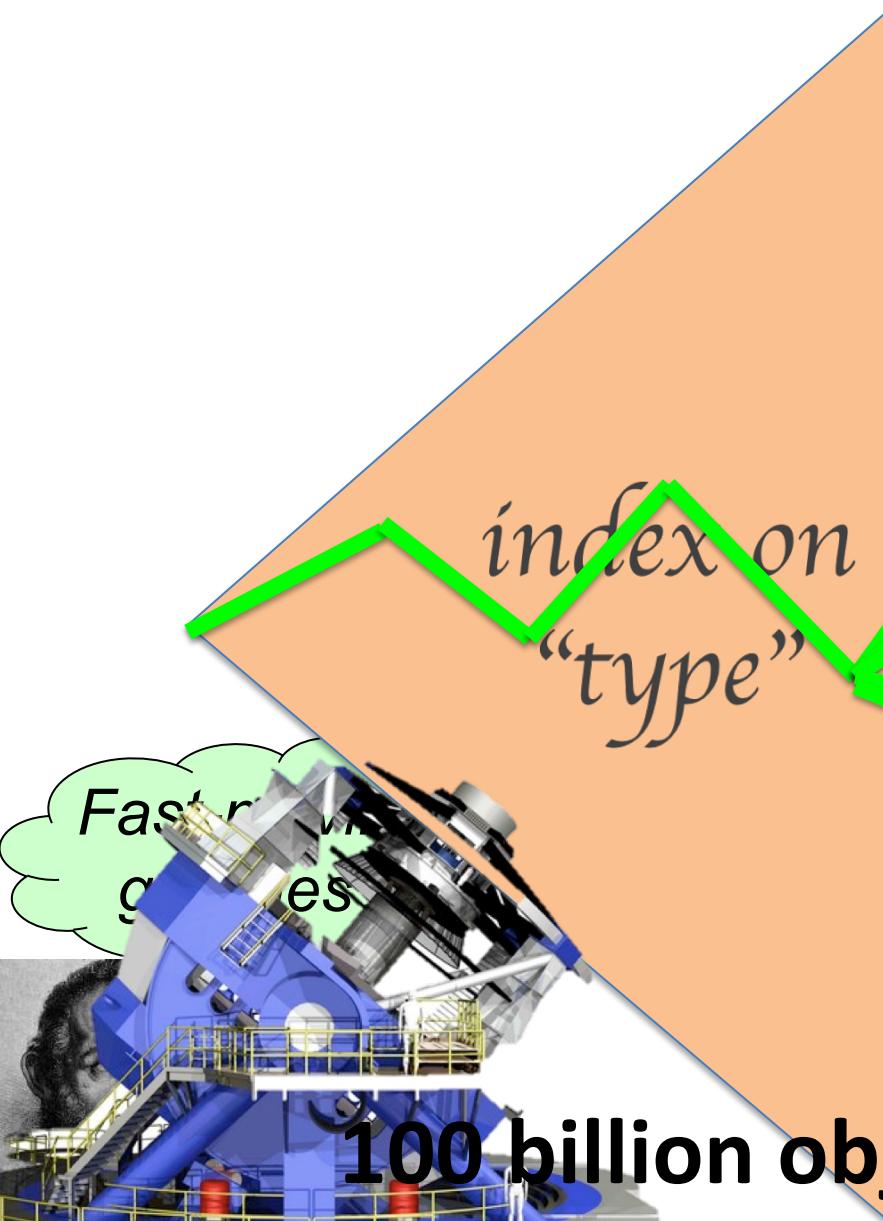
Due 2017

*Find galaxy
clusters*



...in a second?

Large Synoptic Survey Telescope (LSST)



Astronomy DBs

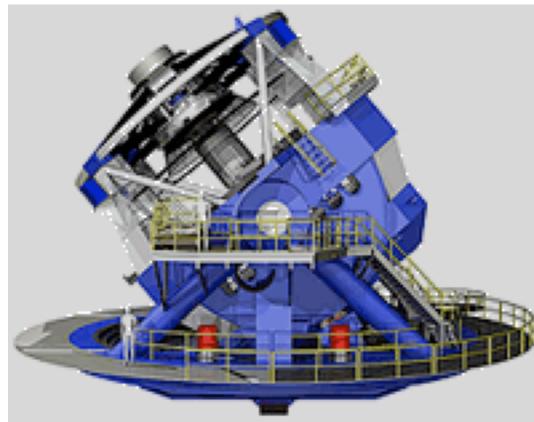


Table with Observations

	TYPE	CX	CY	E1	E2	E3	...

Unused part

Q1

Q2
REPLACE!

	TYPE	CX	CY

	E1	E2	E3

	...

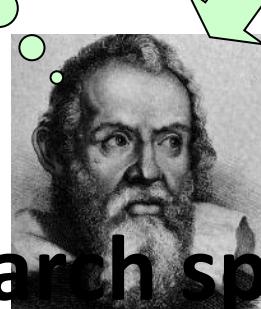
Index

	TYPE	CX	CY

E1

E2	E3

Find galaxy clusters



Search space exponentially growing

Designing Astronomy DBs

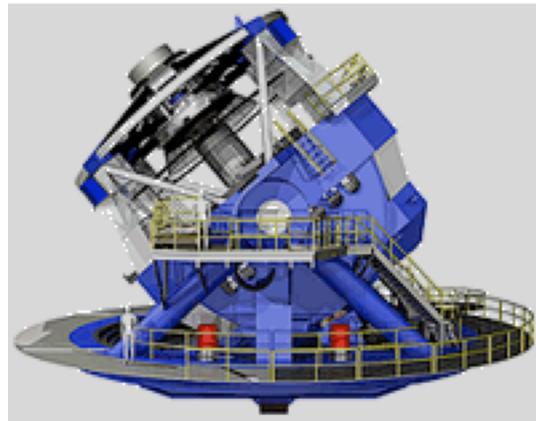
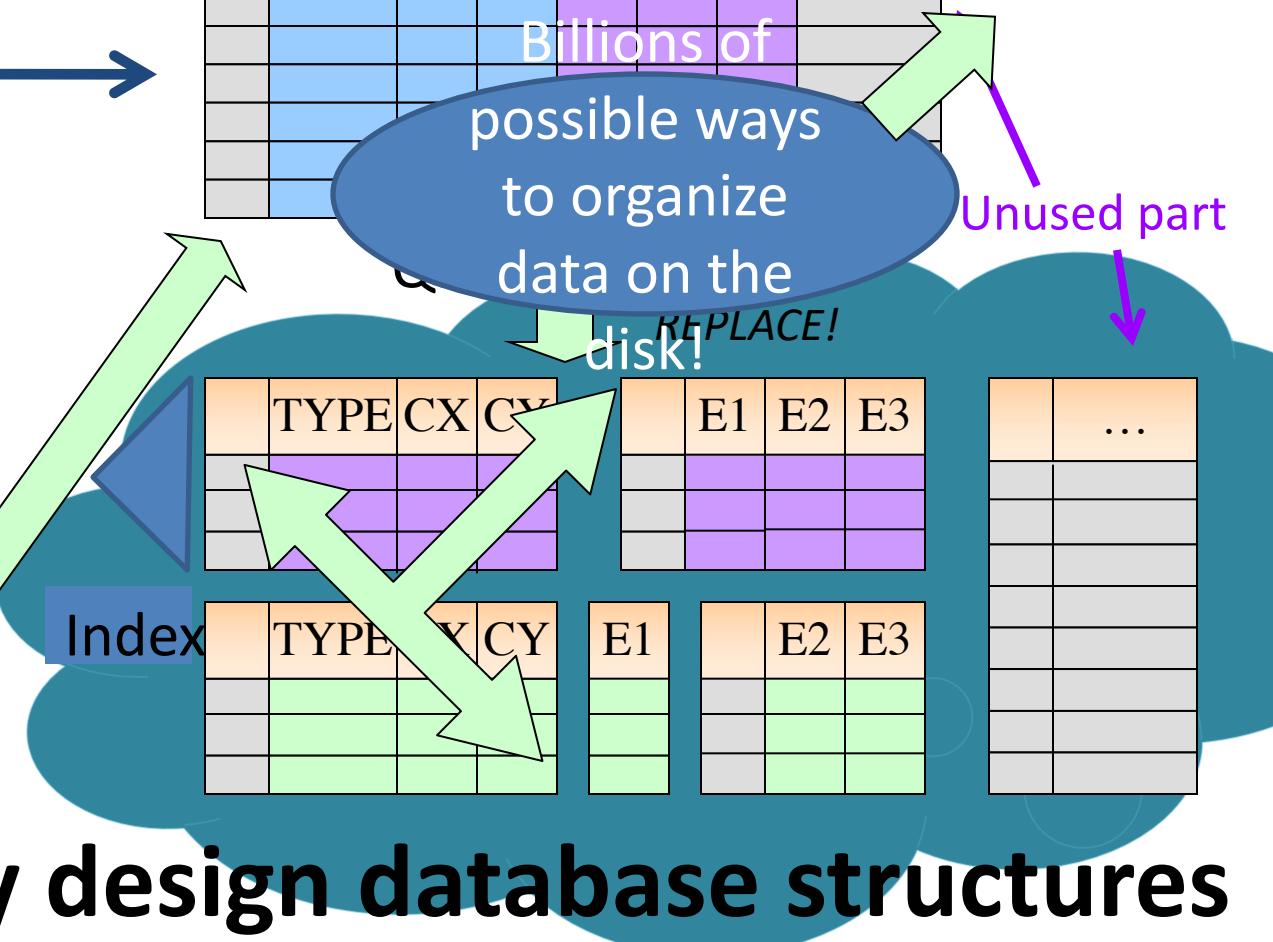
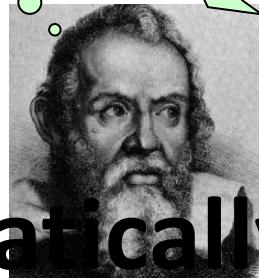


Table with

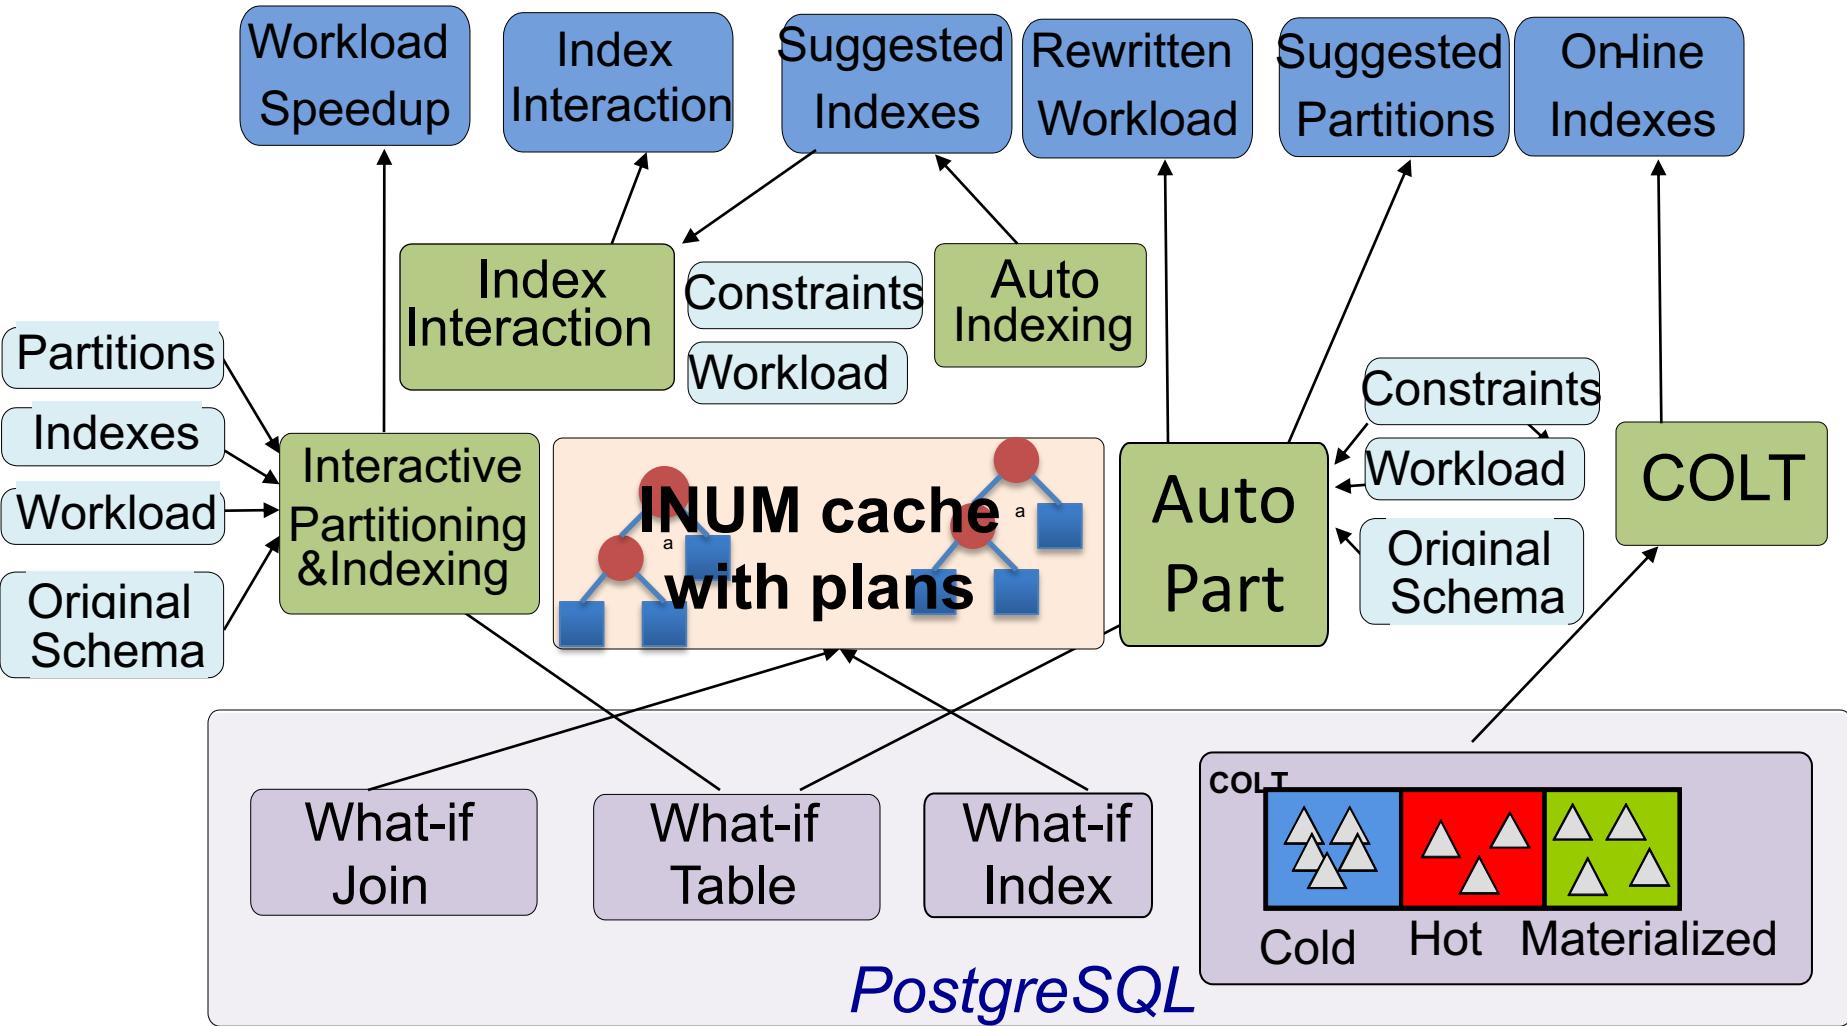
possible ways
to organize
data on the
disk!

Unused part

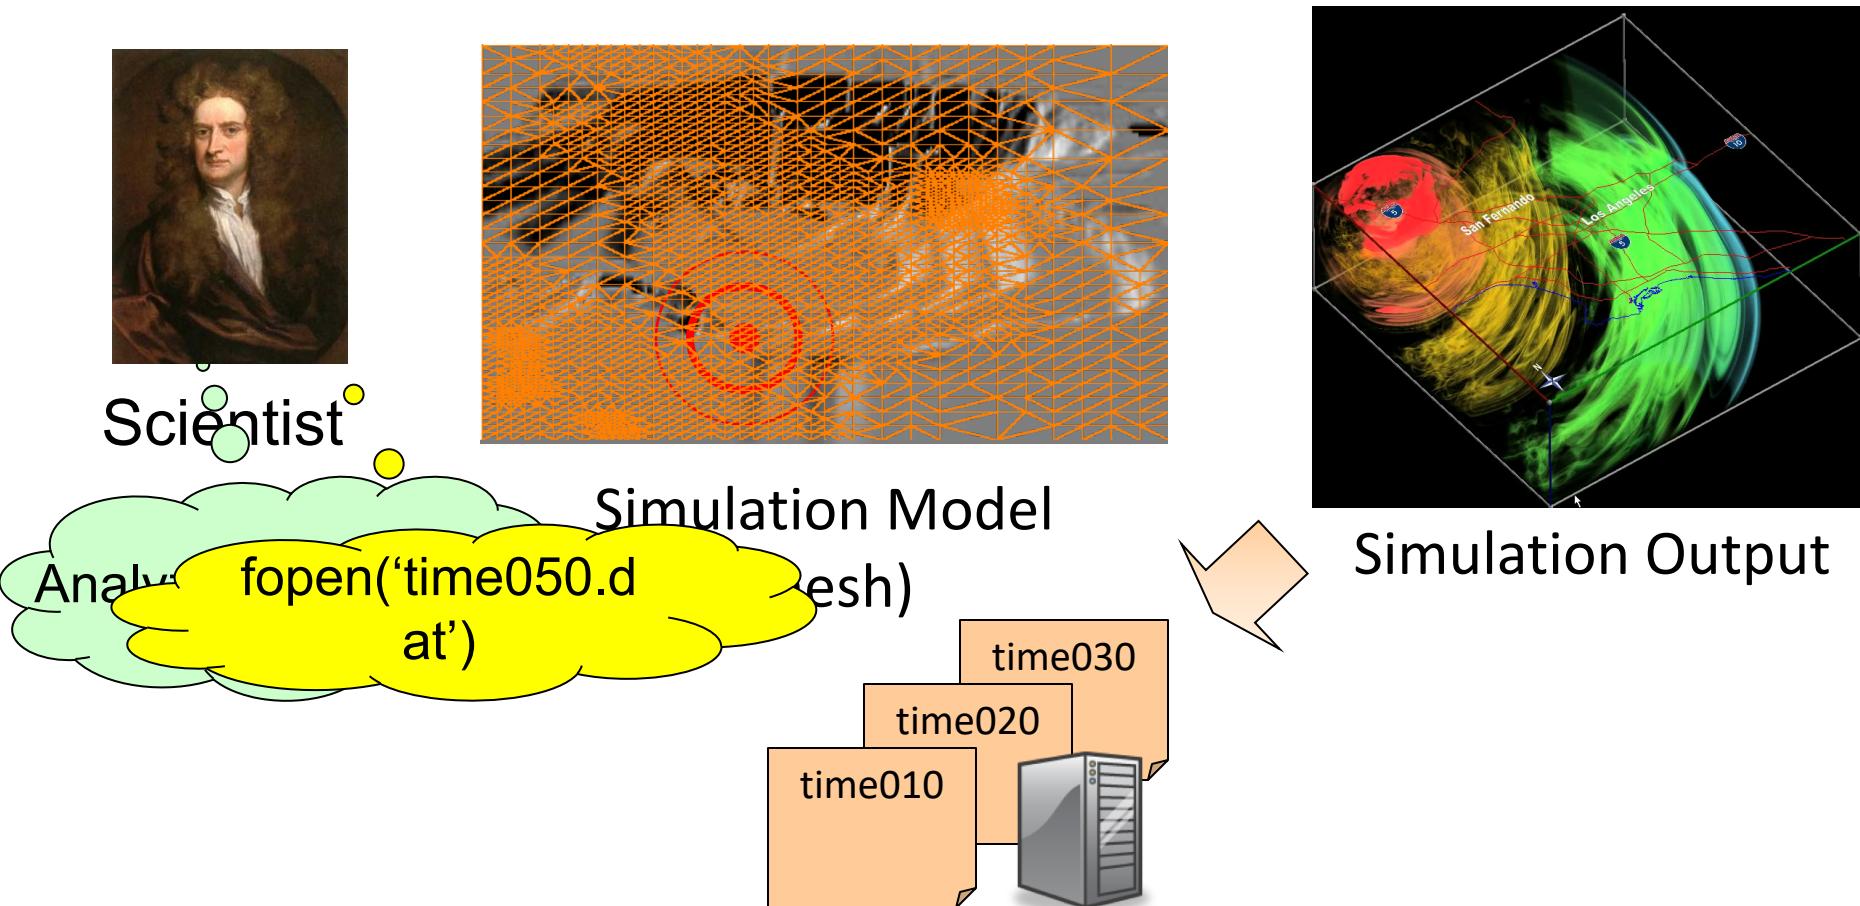
Find galaxy clusters



Database Designer

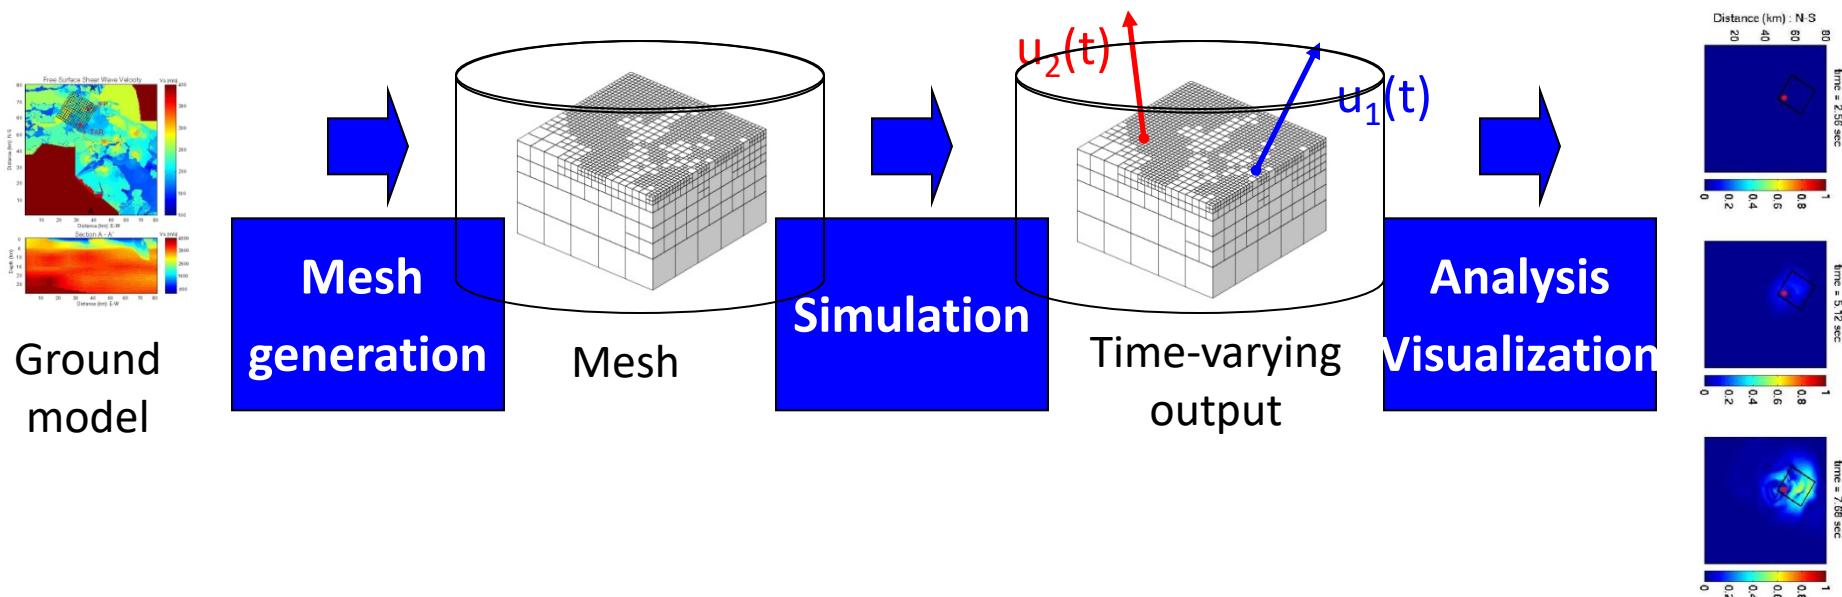


Earthquake Simulation/Analysis

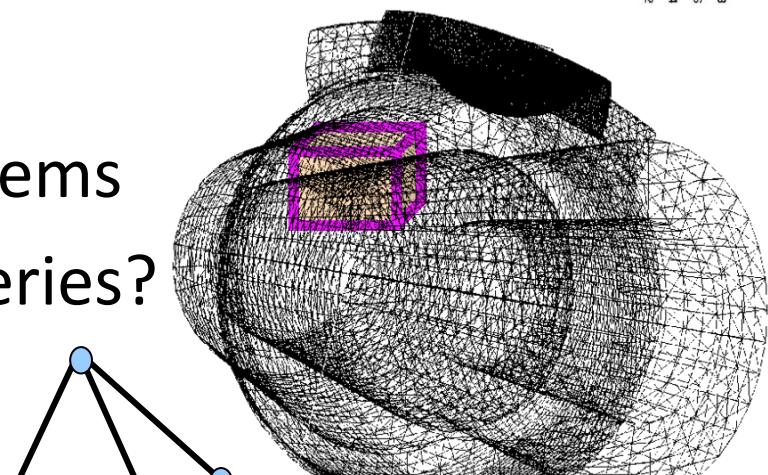


**100GB per time-step, 20000 time-steps
performance? access to complex structures?**

Earthquake Simulation

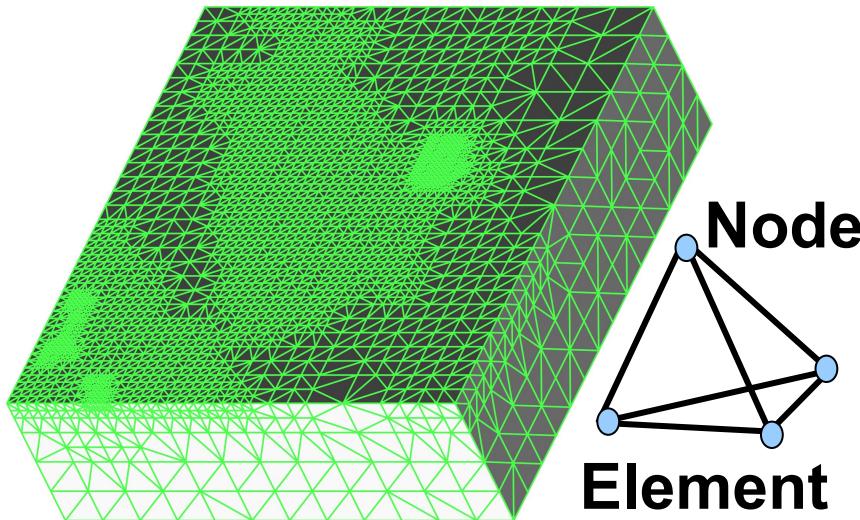


- Huge models, huge outputs
- Not scalable, need database systems
- Mesh representation for fast queries?

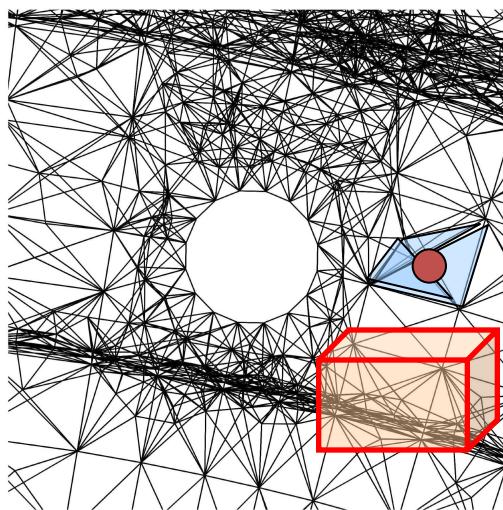


Search entire mesh for regions (not scalable)

Tetrahedral Mesh Models



- Queries
 - Point
 - Range
 - Feature

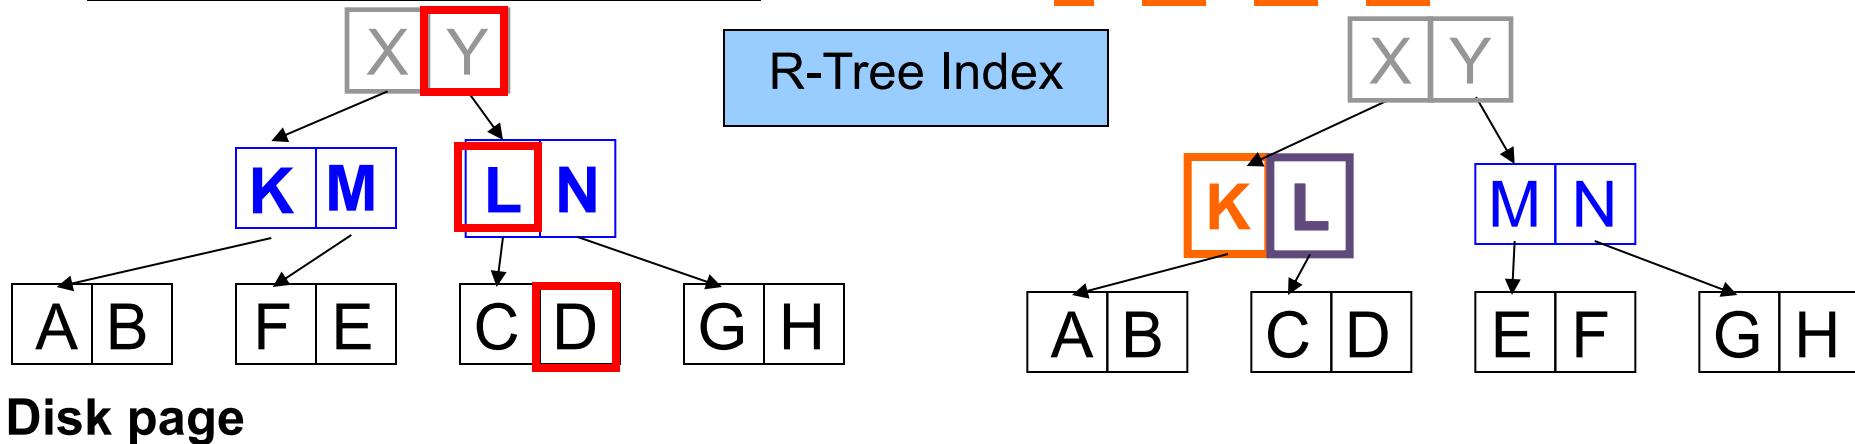
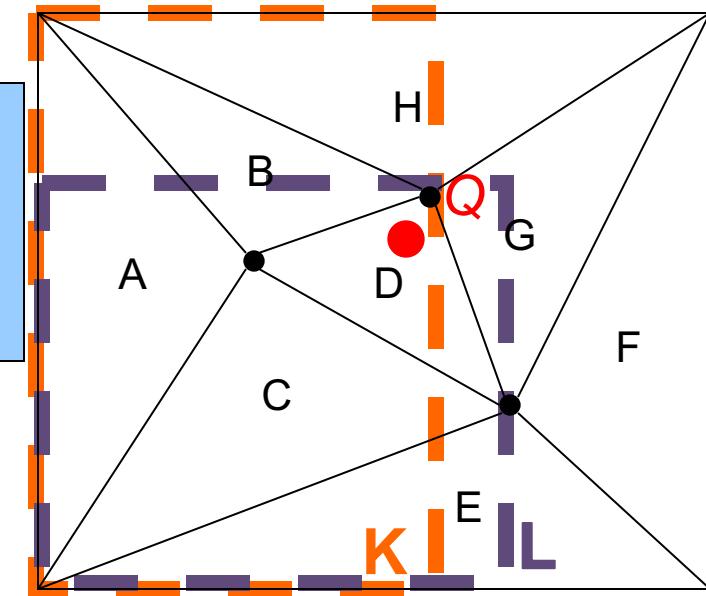
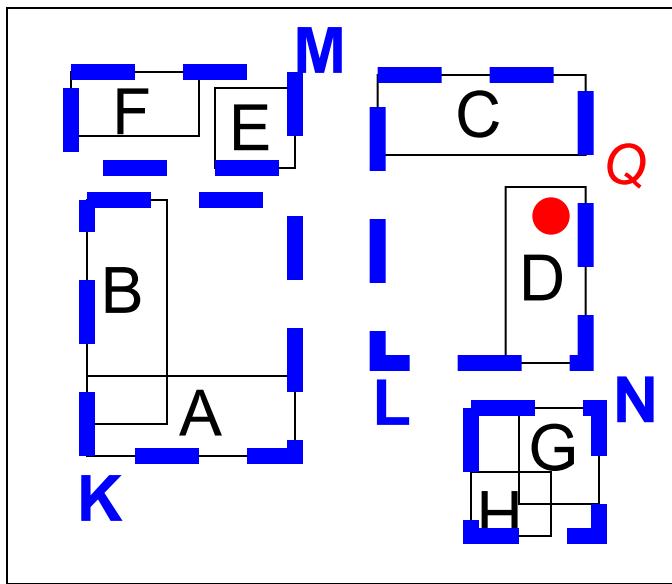


- Example: Visualization
 - Show ground velocity at **Q**
 - Draw the temperature of **R**

Courtesy Cornell Fracture Group

Goal: efficiently process geometric queries

R-Tree Indexing

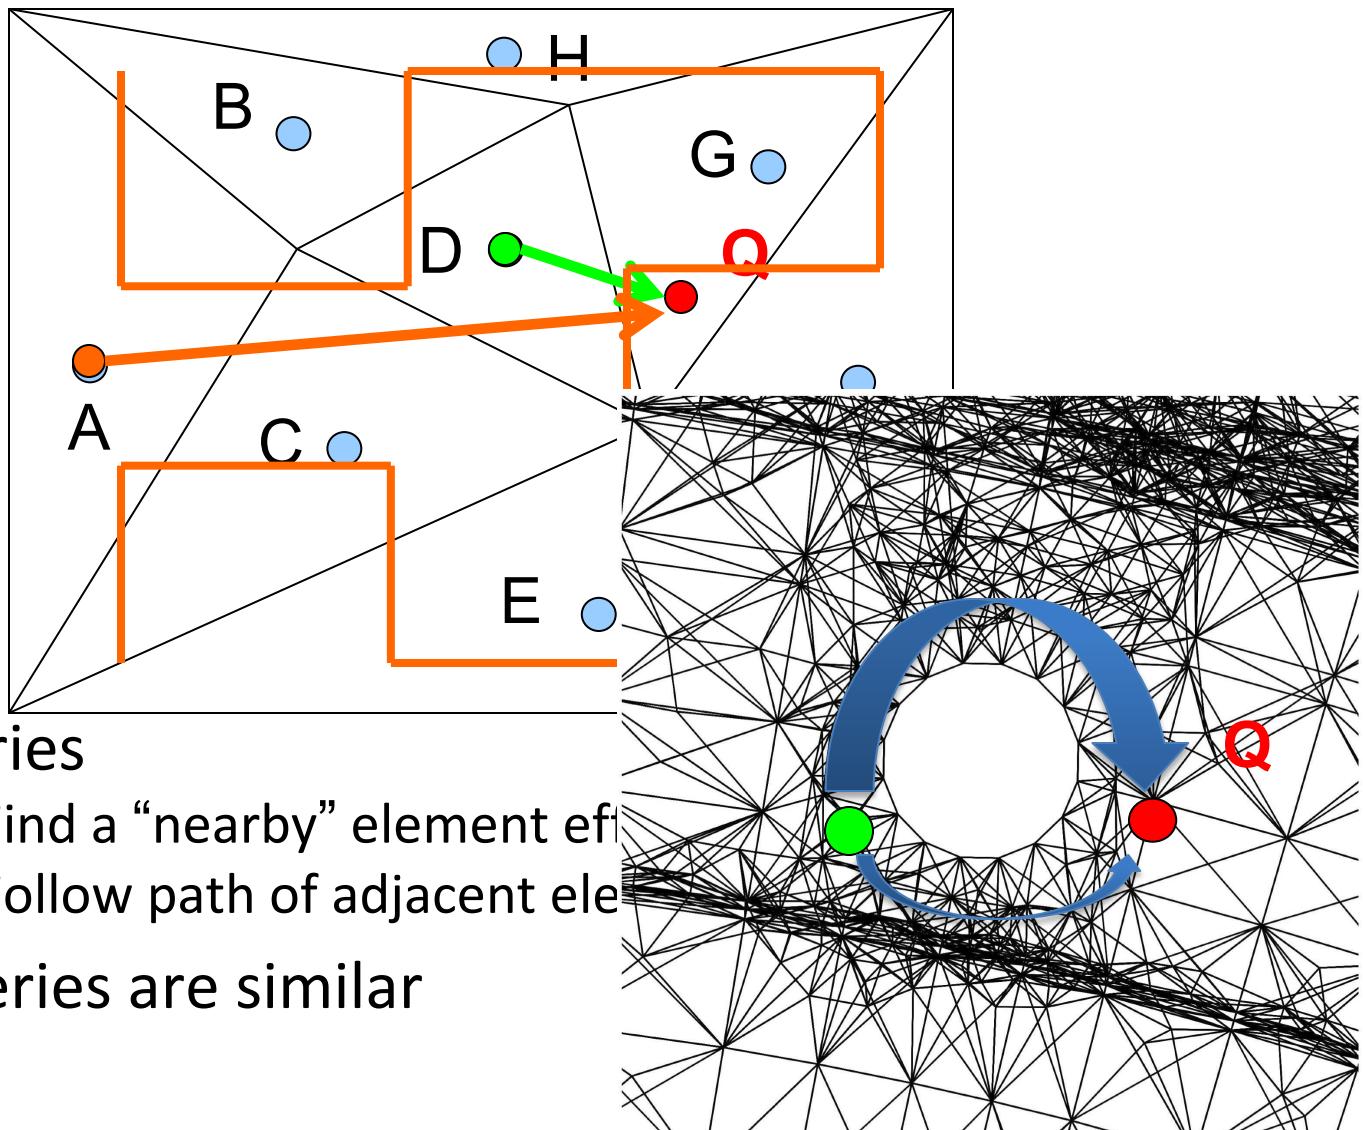


Tight mesh connectivity hurts performance

Solution: Fast Indexing of meshes

- Tetrahedral meshes
 - Index first the centers of the tetrahedra
 - Try to land somewhere near point of interest
 - Use triangulation walking to reach target
- Brain simulation
 - Index an overlay of the brain model
 - Use Voronoi tessellation to fill empty spaces
 - Enable triangulation walking to target

Directed Local Search



- Point queries
 - Step 1: Find a “nearby” element efficiently
 - Step 2: Follow path of adjacent elements
- Range queries are similar

Conclusions

- Data Management has undergone tremendous changes:
 - New requirements
 - New use cases
- Current needs of scientific applications still not served well
 - Spatial data
 - Large scale analysis