

# SQL

## Tutorial II

Valentin CLÉMENT - [vec16@ic.ac.uk](mailto:vec16@ic.ac.uk)

Department of Computing

## Accessing the test server - reminder

You should all have read access to the following PostgreSQL database:

Host : db.doc.ic.ac.uk

Port : 5432

DB : vec16

Username : *Your college username*

Password : Your **Postgres** password, which is **not** your college password! You can find it in your emails as **Your new PostgreSQL password and database for [...]**

In this DB, you only have access to selected *schemas*, such as *sn*.

- ▶ You have a similar DB called 'your username', everything else being similar. You only have very limited access on mine, you should have full access on yours, letting you go free
- ▶ If you want to copy one or more table, the following command might be helpful: `CREATE TABLE staff AS TABLE vec16.doc.staff`: it creates a table as the result of a query, in this case copying it.

- ▶ Today we will tackle more intricate SQL with a small 'social network' example
- ▶ Social Networks are interesting, as they exemplify *graphs*:

### Definition

A graph is an ordered pair  $G = (V, E)$  comprising a set  $V$  of vertices together with a set  $E$  of edges, which are 2-element subsets of  $V$

- ▶ That sounds an awful lot like relations...

- ▶ Today we will tackle more intricate SQL with a small 'social network' example
- ▶ Social Networks are interesting, as they exemplify *graphs*:

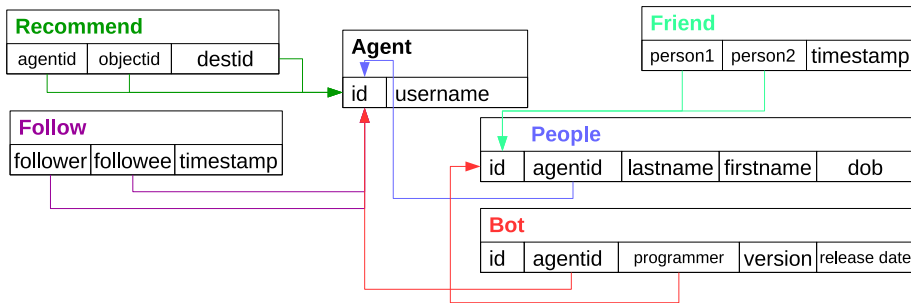
### Definition

A graph is an ordered pair  $G = (V, E)$  comprising a set  $V$  of vertices together with a set  $E$  of edges, which are 2-element subsets of  $V$

- ▶ That sounds an awful lot like relations...
- ▶ And indeed, graphs are fundamentally isomorphic to relations  $vertex(\underline{id}, \dots, vlab), edge(\underline{vertexid}_1, \underline{vertexid}_2, \dots, elab)$

- ▶ The network represents the interaction of **agent**(id, username)
- ▶ Of course, this is the 21<sup>st</sup> century, so agents can be **people**(id, **agentid**, lastname, firstname, dob) or **bot**(id, **agentid**, **programmer**, version, releasedate)
- ▶ An agent can **follow**(**follower**, **followee**, timestamp) another agent, human or not
- ▶ Two humans can be **friend**(**person1**, **person2**, timestamp).
- ▶ Any agent can **recommend**(**agentid**, **objectid**, **destid**) another agent to a third one.

## Representation - *sn* schema



**sn.agent** (id, username)

**sn.people** (id, agentid(agent.id), lastname, firstname, dob)

**sn.bot** (id, agentid(agent.id), programmer(people.id), version, release)

**sn.friend** (person1(people.id), person2(people.id), timestamp)

**sn.recommend** (agentid(agent.id), objectid(agent.id), destid(agent.id))

**sn.follow** (follower(agent.id), followee(agent.id), timestamp)

- ▶ List the first 50 *people* as: agent id, username, first name, last name and date of birth, ordered by *descending* date of birth.
- ▶ List, for each person, their first and last name and the number of bots they have programmed (different versions count as different bots), ordered by descending number of bots

Relevant tables:

**sn.agent** (id, username)

**sn.people** (id, agentid(agent.id), lastname, firstname, dob)

**sn.bot** (id, agentid(agent.id), programmer(people.id), version, release)

- List people by descending DOB:

```
SELECT a.id, a.username, p.firstname,  
        p.lastname, p.dob  
FROM sn.agent a  
JOIN sn.people p on p.agentid=a.id  
ORDER BY p.dob DESC  
LIMIT 50;
```



- List people by descending DOB:

```
SELECT a.id, a.username, p.firstname,  
        p.lastname, p.dob  
FROM sn.agent a  
JOIN sn.people p on p.agentid=a.id  
ORDER BY p.dob DESC  
LIMIT 50;
```

- List people by bot programmed:

```
SELECT p.firstname, p.lastname, count(b.id)  
FROM sn.people as p  
JOIN sn.bot as b ON b.programmer = p.id  
GROUP BY p.id  
ORDER BY count(b.id) DESC;
```

- ▶ List all the people's firstname, lastname and bot username who have programmed a bot which followed them.
- ▶ For each *person*, list their last name, username and the total number of agents that *follow them*
- ▶ For each person, list their last name, username, the total number of people that follow them **and** the total number of agents that *they follow*
- ▶ Some people might have *no* followers and/or follow no one: make sure you still list them !

Relevant tables:

**sn.agent** (id, username)

**sn.people** (id, agentid(agent.id), lastname, firstname, dob)

**sn.follow** (follower(agent.id), followee(agent.id), timestamp)

- ▶ List all the people's firstname, lastname and bot username who have programmed a bot which followed them.
- ▶ For each *person*, list their last name, username and the total number of agents that *follow them*
- ▶ For each person, list their last name, username, the total number of people that follow them **and** the total number of agents that *they follow*
- ▶ Some people might have *no* followers and/or follow no one: make sure you still list them ! Hint: This means **[INNER] JOIN** will not be enough

Relevant tables:

**sn.agent** (id, username)

**sn.people** (id, agentid(agent.id), lastname, firstname, dob)

**sn.follow** (follower(agent.id), followee(agent.id), timestamp)

- Users programming bots to follow themselves:

```
SELECT p.firstname, p.lastname, a.username  
FROM sn.people p  
JOIN sn.bot b ON b.programmer = p.id  
JOIN sn.follow f ON f.follower = b.agentid  
JOIN sn.agent a ON b.agentid=a.id  
WHERE f.followee = p.agentid;
```

- ▶ Number of followers:

```
SELECT p.lastname, a.username,  
        count(f.id) AS followers  
FROM sn.agent a  
JOIN sn.people p ON p.agentid = a.id  
LEFT JOIN sn.follow f ON f.followee = a.id  
GROUP BY a.id, a.username, p.lastname  
ORDER BY a.id;
```

- ▶ Note the **LEFT JOIN**, necessary to account for people with no followers/following
- ▶ We **GROUP BY** the id and use it for ordering out of convenience: username/lastname are *not guaranteed* to be unique, whereas ids will provide us reliable and reproducible results.

- Followers and following:

```
SELECT p.lastname, a.username,  
        count(distinct f1.id) as followers,  
        count(distinct f2.id) as following  
FROM sn.agent a  
JOIN sn.people p ON p.agentid = a.id  
LEFT JOIN sn.follow f1 ON f1.followee = a.id  
LEFT JOIN sn.follow f2 ON f2.follower = a.id  
GROUP BY a.id, a.username, p.lastname  
ORDER BY a.id;
```

- Approach is very similar to the previous one
- Note we now need to count **distinct** ids: this is because the two full joins return all valid pairs

- ▶ List all the friends of John Doe, in format lastname, firstname.
- ▶ Remember: friend relation can go either way: either  $A$  is a friend of  $B$ , or  $B$  is a friend of  $A$
- ▶ There are duplicates in the database

Relevant tables:

`sn.people` (id, agentid(agent.id), lastname, firstname, dob)

`sn.friend` (person1(people.id), person2(people.id), timestamp)

- Find the ID of John Doe:

```
SELECT p.id FROM sn.people p
WHERE p.lastname = 'Doe' AND p.firstname = 'John';
```

- Find his friends:

```
SELECT p.lastname, p.firstname
FROM sn.people p
WHERE p.id IN (
    SELECT DISTINCT p.id
    FROM sn.people p
    WHERE p.id IN (
        SELECT f.person1 FROM sn.friend f WHERE f.person2 = 1
    ) OR p.id IN (
        SELECT f.person2 FROM sn.friend f WHERE f.person1 = 1
    )
);
```



In all its horrible (but side-effect free !) glory

- List all friends of John Doe

```
SELECT p.lastname, p.firstname FROM sn.people p WHERE p.id IN (  
  SELECT DISTINCT p.id FROM sn.people p WHERE p.id IN (  
    SELECT p.id FROM sn.people p  
    JOIN sn.friend f on p.id = f.person1  
    JOIN sn.people q on q.id = f.person2  
    WHERE q.id IN (  
      SELECT p.id FROM sn.people p  
      WHERE p.lastname = 'Doe' AND p.firstname = 'John'  
    )  
  ) OR p.id IN (  
    SELECT p.id FROM sn.people p  
    JOIN sn.friend f on p.id = f.person2  
    JOIN sn.people q on q.id = f.person1  
    WHERE q.id IN (  
      SELECT p.id FROM sn.people p  
      WHERE p.lastname = 'Doe' AND p.firstname = 'John'  
    )  
  )  
);
```

```
WITH jd(id) AS (  
  SELECT p.id FROM sn.people p WHERE p.lastname = 'Doe' AND p.firstname = 'John'  
) , friendsofjohndoe(id) AS (  
  SELECT DISTINCT p.id FROM sn.people p  
  WHERE p.id IN (  
    SELECT f.person1 FROM sn.friend f WHERE f.person2 IN (SELECT id FROM jd)  
  ) OR p.id IN (  
    SELECT f.person2 FROM sn.friend f WHERE f.person1 IN (SELECT id FROM jd)  
  )  
)  
SELECT p.lastname, p.firstname FROM sn.people p  
JOIN friendsofjohndoe fojd ON p.id = fojd.id;
```

- ▶ For each *person* who has been the destination of a recommendation at least once, list all recommending agents' username, along with their number of recommendations
- ▶ Same as above, but limit only to recommendations by other *people*
- ▶ Advanced level: Same as above, but limit only to recommendation by their *friends*
- ▶ SQL Grandmaster level: for each *person*, find a single *person*, who is not their friend, and is most recommended by their friends

Relevant tables:

`sn.agent` (id, username)

`sn.people` (id, agentid(agent.id), lastname, firstname, dob)

`sn.recommend` (agentid(agent.id), objectid(agent.id), destid(agent.id))

`sn.friend` (person1(people.id), person2(people.id), timestamp)

- ▶ Let's unpack the question: 'for each person [...], list all [...] username'

- ▶ Let's unpack the question: 'for each person [...], list all [...] username'
- ▶ This means generating a pairs (person, username), linking the person and the recommender, something like:

```
SELECT p.lastname, p.firstname, a.username  
FROM sn.people p  
JOIN sn.agent a ON ???  
WHERE ???
```

- ▶ Let's unpack the question: 'for each person [...], list all [...] username'
- ▶ This means generating a pairs (person, username), linking the person and the recommender, something like:

```
SELECT p.lastname, p.firstname, a.username  
FROM sn.people p  
JOIN sn.agent a ON ???  
WHERE ???
```

- ▶ '[...] along with their number of recommendations':

```
SELECT p.lastname, p.firstname, a.username, COUNT(r.id)  
FROM sn.people p  
JOIN sn.agent a ON ???  
JOIN sn.recommend r ON ???  
WHERE ???  
GROUP BY p.id, p.lastname, p.firstname, a.id, a.username
```

- All that remains is the joining to perform. . . This is driven by the *recommendations*, with the person being the recipient and the agent the sender:

```
SELECT p.lastname, p.firstname, a.username, COUNT(r.id)
FROM sn.people p
JOIN sn.recommend r ON r.destid = p.agentid
JOIN sn.agent a ON r.agentid = a.id
WHERE ???
GROUP BY p.id, p.lastname, p.firstname, a.id, a.username
```

- All that remains is the joining to perform. . . This is driven by the *recommendations*, with the person being the recipient and the agent the sender:

```
SELECT p.lastname, p.firstname, a.username, COUNT(r.id)
FROM sn.people p
JOIN sn.recommend r ON r.destid = p.agentid
JOIN sn.agent a ON r.agentid = a.id
WHERE ???
GROUP BY p.id, p.lastname, p.firstname, a.id, a.username
```

- Actually, no filtering (WHERE) is needed! So we get the final query:

```
SELECT p.lastname, p.firstname, a.username, COUNT(r.id)
FROM sn.people p
JOIN sn.recommend r ON r.destid = p.agentid
JOIN sn.agent a ON r.agentid = a.id
GROUP BY p.id, p.lastname, p.firstname, a.id, a.username
```



- ▶ Same as above, but limit only to recommendations by other *people*

- ▶ Same as above, but limit only to recommendations by other *people*
- ▶ Option 1: simply add a filter:

```
SELECT p.lastname, p.firstname, a.username, COUNT(r.id)
FROM sn.people p
JOIN sn.recommend r ON r.destid = p.agentid
JOIN sn.agent a ON r.agentid = a.id
WHERE a.id IN (SELECT agentid FROM sn.people)
GROUP BY p.id, p.lastname, p.firstname, a.id, a.username
```

- ▶ Same as above, but limit only to recommendations by other *people*
- ▶ Option 1: simply add a filter:

```
SELECT p.lastname, p.firstname, a.username, COUNT(r.id)
FROM sn.people p
JOIN sn.recommend r ON r.destid = p.agentid
JOIN sn.agent a ON r.agentid = a.id
WHERE a.id IN (SELECT agentid FROM sn.people)
GROUP BY p.id, p.lastname, p.firstname, a.id, a.username
```

- ▶ Option 2: join

```
SELECT p.lastname, p.firstname, a.username, COUNT(r.id)
FROM sn.people p
JOIN sn.recommend r ON r.destid = p.agentid
JOIN sn.agent a ON r.agentid = a.id
JOIN sn.people q ON a.id = q.agentid
GROUP BY p.id, p.lastname, p.firstname, a.id, a.username
```

- ▶ Advanced level: same as above, but limit to recommendations by that person's *friends*

- ▶ Advanced level: same as above, but limit to recommendations by that person's *friends*
- ▶ Need to join in the *friend*(person, person) many-to-many relationship. . .

- ▶ Advanced level: same as above, but limit to recommendations by that person's *friends*
- ▶ Need to join in the *friend*(person, person) many-to-many relationship. . .
- ▶ Using the second option from above, that is actually straightforward:

```
SELECT p.lastname, p.firstname, a.username, COUNT(r.id)
FROM sn.people p
JOIN sn.recommend r ON r.destid = p.agentid
JOIN sn.agent a ON r.agentid = a.id
JOIN sn.people q ON a.id = q.agentid
JOIN sn.friend f ON f.person1 = p.id
WHERE f.person2 = q.id
GROUP BY p.id, p.lastname, p.firstname, a.id, a.username
```

- ▶ Or is it. . . ?

- ▶ Advanced level: same as above, but limit to recommendations by that person's *friends*
- ▶ Need to join in the *friend*(person, person) many-to-many relationship. . .
- ▶ Using the second option from above, that is actually straightforward:

```
SELECT p.lastname, p.firstname, a.username, COUNT(r.id)
FROM sn.people p
JOIN sn.recommend r ON r.destid = p.agentid
JOIN sn.agent a ON r.agentid = a.id
JOIN sn.people q ON a.id = q.agentid
JOIN sn.friend f ON f.person1 = p.id
WHERE f.person2 = q.id
GROUP BY p.id, p.lastname, p.firstname, a.id, a.username
```

- ▶ Or is it. . . ?
- ▶ Try inverting 'person1' and 'person2' in the query above. . .
- ▶ The non-oriented relationship raises its ugly head again!

- ▶ Doing a big WITH extravaganza as in previous example is possible...



- ▶ Doing a big WITH extravaganza as in previous example is possible...
- ▶ But there is another solution for these kind of problems: tuple comparison:

```
SELECT p.lastname, p.firstname, a.username, COUNT(r.id)
FROM sn.people p
JOIN sn.recommend r ON r.destid = p.agentid
JOIN sn.agent a ON r.agentid = a.id
JOIN sn.people q ON a.id = q.agentid
WHERE (
    (p.id, q.id) IN (SELECT f.person1, f.person2 FROM sn.friend f) OR
    (p.id, q.id) IN (SELECT f.person2, f.person1 FROM sn.friend f)
)
GROUP BY p.id, p.lastname, p.firstname, a.id, a.username
```

- ▶ Grandmaster level: for each *person*, find the *person*, who is not their friend, and is most recommended by their friends

- ▶ Grandmaster level: for each *person*, find the *person*, who is not their friend, and is most recommended by their friends
- ▶ Proceeding methodically: WHAT are we trying to find ?

```
SELECT p.firstname, p.lastname, q.firstname, q.lastname  
FROM sn.people p  
JOIN sn.people q ON ???  
WHERE ???
```

- Let's make our life simpler first...

```
CREATE TABLE allfriends(p1, p2, i1, i2) AS (  
  SELECT f.person1, f.person2, p.agentid, q.agentid  
  FROM sn.friend f  
  JOIN sn.people p ON p.id = f.person1  
  JOIN sn.people q ON q.id = f.person2  
);  
INSERT INTO allfriends(p1, p2, i1, i2) (  
  SELECT f.person2, f.person1, q.agentid, p.agentid  
  FROM sn.friend f  
  JOIN sn.people p ON p.id = f.person1  
  JOIN sn.people q ON q.id = f.person2  
);  
SELECT p.firstname, p.lastname, q.firstname, q.lastname  
FROM sn.people p, sn.people q  
WHERE (p.id, q.id) NOT IN (SELECT p1, p2 FROM allfriends)
```

- Let's start putting in the recommendations

```
SELECT p.firstname person, p.lastname,  
        q.firstname recommendation, q.lastname, COUNT(r.id)  
FROM sn.people p, sn.people q  
JOIN sn.recommend r ON r.objectid = q.agentid  
WHERE (p.id, q.id) NOT IN (SELECT p1, p2 FROM allfriends) AND  
        (p.id, r.agentid) IN (SELECT p1, i2 FROM allfriends)  
GROUP BY (p.firstname, p.lastname, q.firstname, q.lastname)
```

- ▶ Let's start putting in the recommendations

```
SELECT p.firstname person, p.lastname,  
       q.firstname recommendation, q.lastname, COUNT(r.id)  
FROM sn.people p, sn.people q  
JOIN sn.recommend r ON r.objectid = q.agentid  
WHERE (p.id, q.id) NOT IN (SELECT p1, p2 FROM allfriends) AND  
      (p.id, r.agentid) IN (SELECT p1, i2 FROM allfriends)  
GROUP BY (p.firstname, p.lastname, q.firstname, q.lastname)
```

- ▶ Looks good! But all these aggregations sound really tricky to put in now. . . It is time for a step back.

- ▶ There is an obvious FD from people's id to their names, so let's only get that
- ▶ We will be able to sort ourselves once we have all the good ids.

```
SELECT p.id person, q.id recommendation, COUNT(r.id) num
FROM sn.people p, sn.people q
JOIN sn.recommend r ON r.objectid = q.agentid
WHERE (p.id, q.id) NOT IN (SELECT p1, p2 FROM allfriends) AND
      (p.id, r.agentid) IN (SELECT p1, i2 FROM allfriends)
GROUP BY (p.id, q.id)
```

- ▶ Actually, having numerical IDs, we can now easily get a representative of each number of recommendations:

```
SELECT person, MAX(recommendation), num FROM (  
  SELECT p.id person, q.id recommendation, COUNT(r.id) num  
  FROM sn.people p, sn.people q  
  JOIN sn.recommend r ON r.objectid = q.agentid  
  WHERE (p.id, q.id) NOT IN (SELECT p1, p2 FROM allfriends) AND  
    (p.id, r.agentid) IN (SELECT p1, i2 FROM allfriends)  
  GROUP BY (p.id, q.id)  
) AS tmp  
GROUP BY person, num
```

- ▶ Just have to add something like:

```
WHERE num = maxnum
```

But how ?



- We want *person*, *any(recommendation)*, *num* and the corresponding *person*, *max(num)*...

- ▶ We want *person*, *any(recommendation)*, *num* and the corresponding *person*, *max(num)*...
- ▶ We can generate both, and then JOIN them !

```
WITH tmp(person, rec, num) AS (  
  SELECT p.id person, q.id recommendation, COUNT(r.id) num  
  FROM sn.people p, sn.people q  
  JOIN sn.recommend r ON r.objectid = q.agentid  
  WHERE (p.id, q.id) NOT IN (SELECT p1, p2 FROM allfriends) AND  
         (p.id, r.agentid) IN (SELECT p1, i2 FROM allfriends)  
  GROUP BY (p.id, q.id)  
, r(p, q, num) AS (  
  SELECT person, max(rec), num FROM tmp GROUP BY person, num  
, m(p, maxnum) AS (  
  SELECT person, max(num) FROM tmp GROUP BY person  
)  
SELECT r.p, r.q  
FROM r  
JOIN m ON r.p = m.p  
WHERE r.num = m.maxnum
```

- Let's finally put it back together:

```
WITH tmp(person, rec, num) AS (  
  SELECT p.id person, q.id recommendation, COUNT(r.id) num  
  FROM sn.people p, sn.people q  
  JOIN sn.recommend r ON r.objectid = q.agentid  
  WHERE (p.id, q.id) NOT IN (SELECT p1, p2 FROM allfriends) AND  
    (p.id, r.agentid) IN (SELECT p1, i2 FROM allfriends)  
  GROUP BY (p.id, q.id)  
) , r(p, q, num) AS (  
  SELECT person, max(rec), num FROM tmp GROUP BY person, num  
) , m(p, maxnum) AS (  
  SELECT person, max(num) FROM tmp GROUP BY person  
) , top_rec(p, q) AS (  
  SELECT r.p, r.q FROM r JOIN m ON r.p = m.p WHERE r.num = m.maxnum  
)  
SELECT p.firstname, p.lastname, q.firstname rec_first, q.lastname rec_last  
FROM sn.people p JOIN top_rec r ON r.p = p.id JOIN sn.people q ON r.q = q.id  
ORDER BY p.lastname, p.firstname
```

- ▶ You won't have any monster like that on the exam
- ▶ But this should show you that **SQL is very powerful** (and quite rich)
- ▶ Several sites (most notably Stack Overflow) run on pretty much nothing more than SQL
- ▶ Train, experiment and build your intuition! These tutorials give you examples and syntax, but only practice makes perfect
- ▶ **Use your DoC database**; most languages have fairly straightforward PostgreSQL bindings, which lets you use your DoC database to as a globally accessible store for any project you have
- ▶ Build your queries in steps - what do you need; what combination of sources will get you that; what then needs to be filtered. And finally, in what order do you want the results.

**Questions ?**